

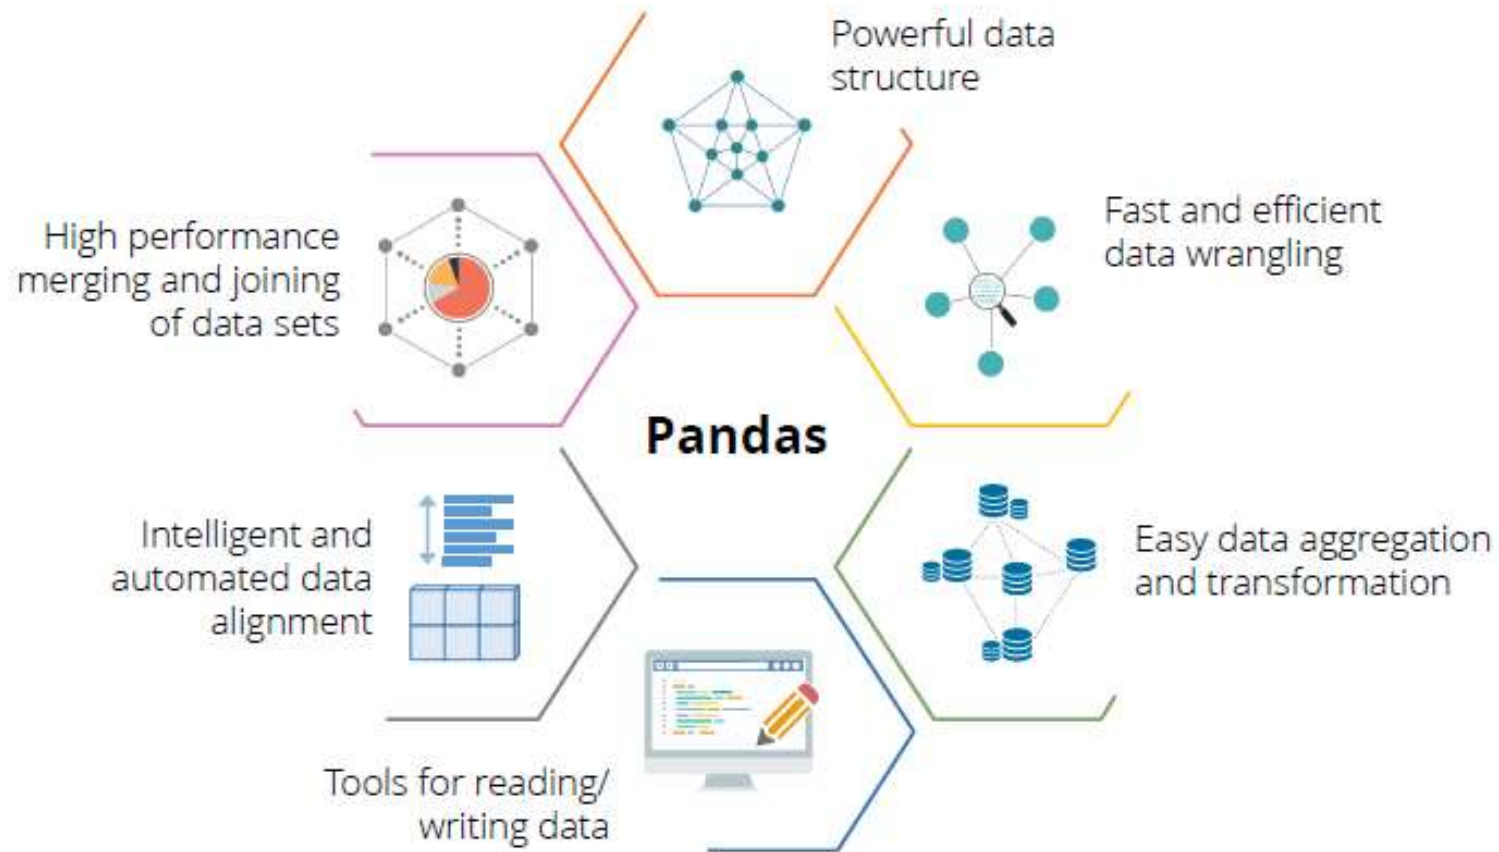
Pandas



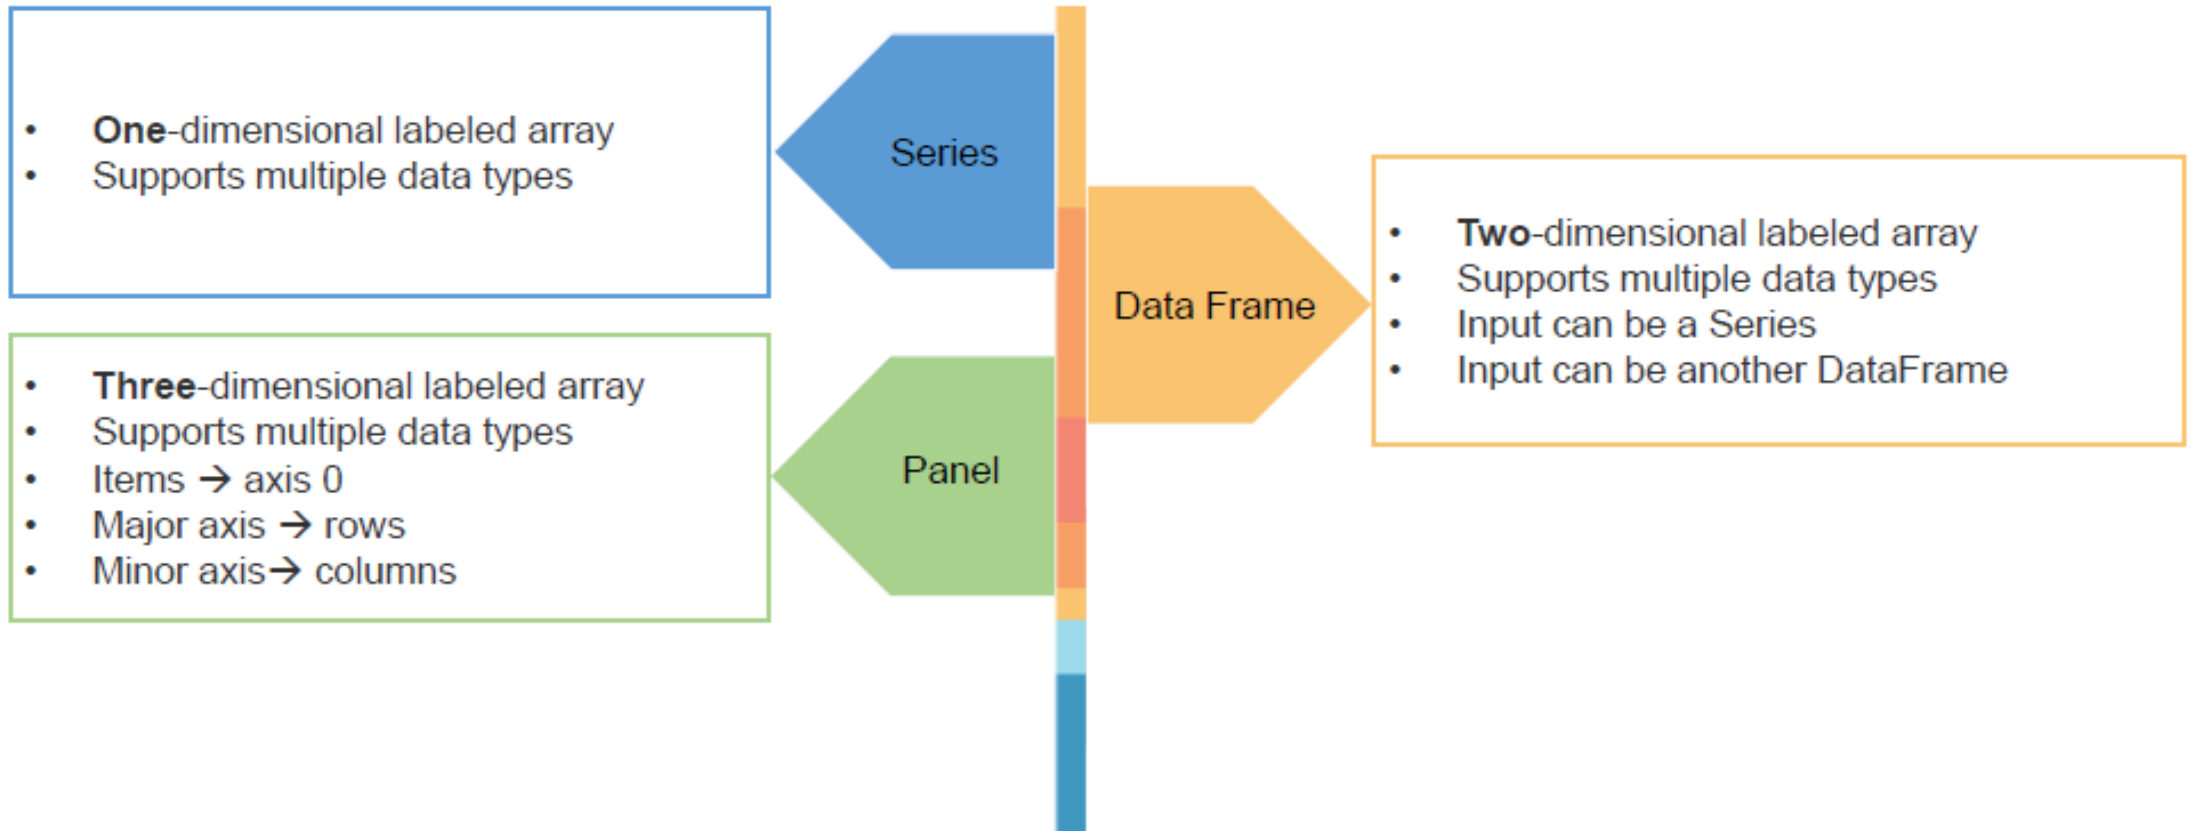
Data Manipulation with Pandas

Pandas Features

The various features of Pandas makes it an efficient library for Data Scientists.




The four main libraries of Pandas data structure are:



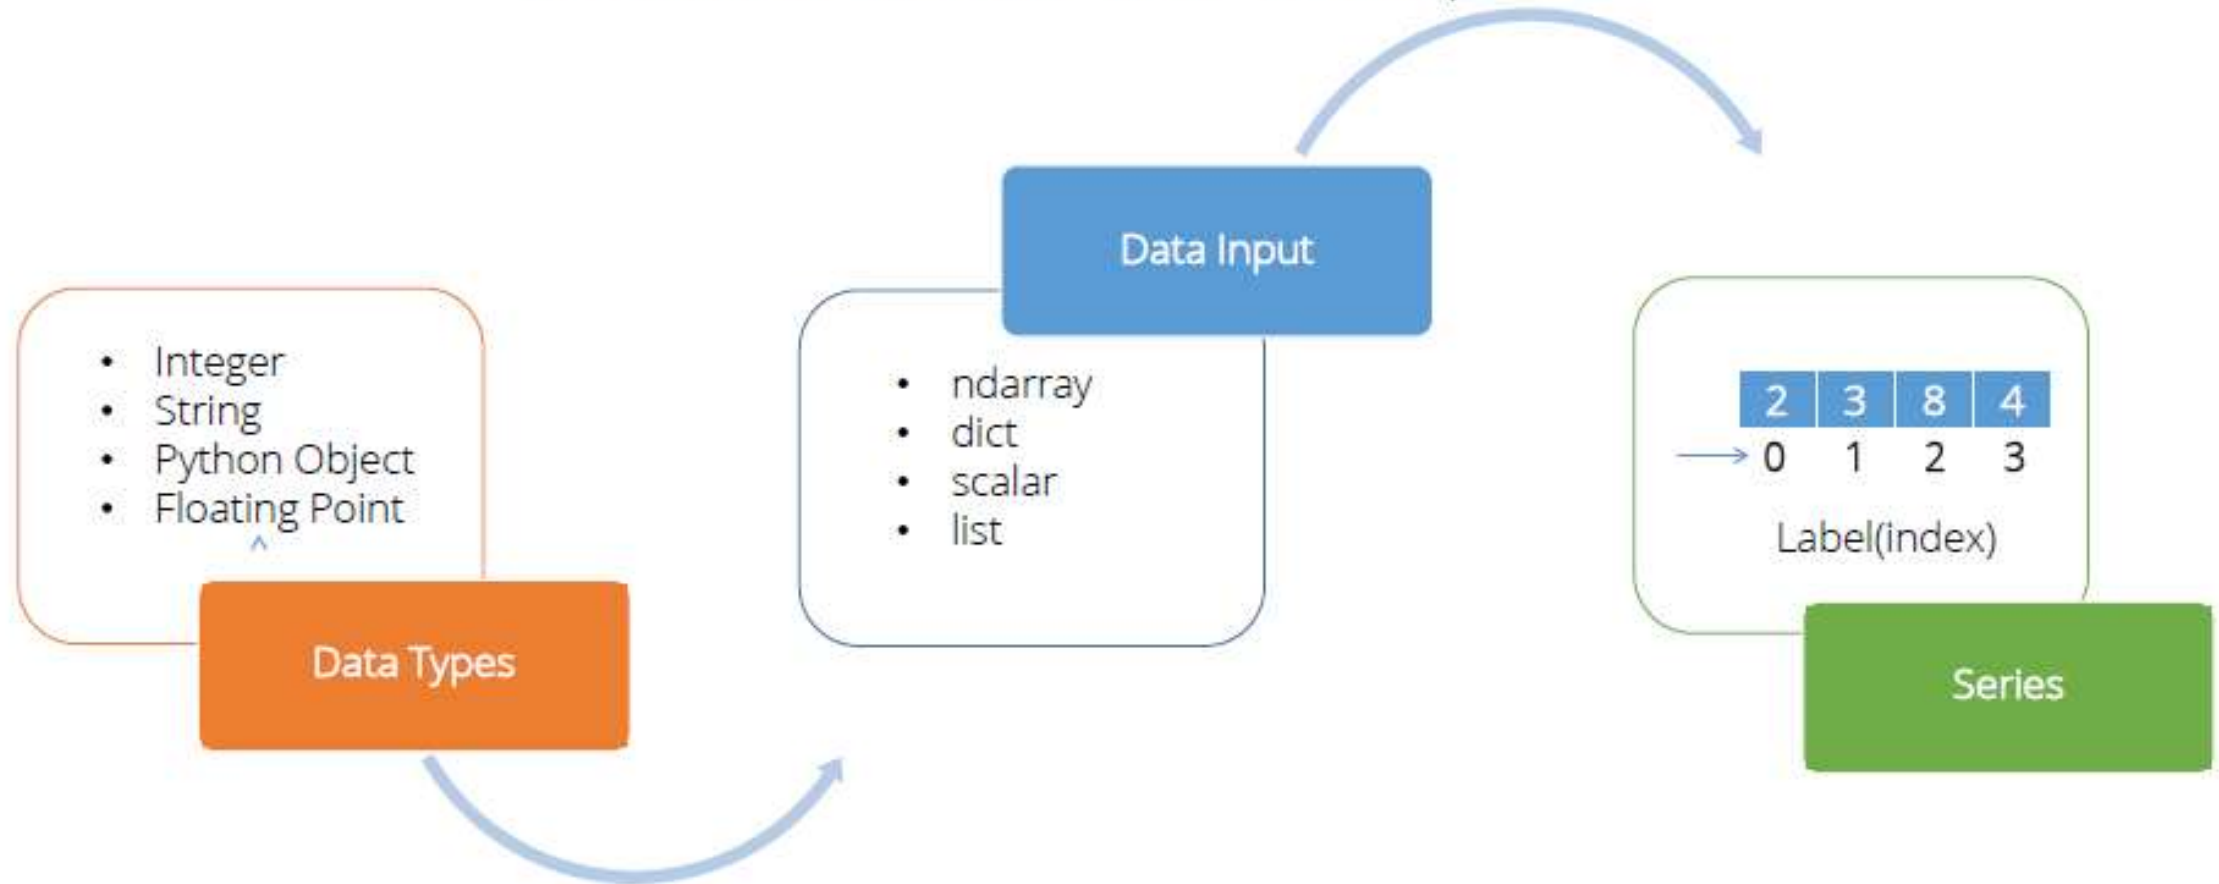
Understanding Series

Series is a one-dimensional array-like object containing data and labels (or index).



Data alignment is intrinsic and will not be broken until changed explicitly by program.

Series can be created with different data inputs:



How to Create Series

Key points to note while creating a series are as follows:

- Import Pandas as it is the main library
- Apply the syntax and pass the data elements as arguments
- Import NumPy while working with ndarrays

Basic Method

```
S = pd.Series(data, index = [index])
```



4	11	21	36
---	----	----	----

Series

Create Series from List

This example shows you how to create a series from a list:

```
In [14]: import numpy as np  
import pandas as pd
```

← Import libraries

```
In [15]: first_series = pd.Series(list('abcdef'))
```

← Pass list as an argument

```
In [16]: print (first_series)
```

0	a
1	b
2	c
3	d
4	e
5	f

Index →

← Data value

dtype: object ← Data type



We have not created index for data but notice that data alignment is done automatically

Create Series from ndarray

This example shows you how to create a series from an ndarray:

```
In [17]: np_country = np.array(['Luxembourg', 'Norway', 'Japan', 'Switzerland', 'United States', 'Qatar', 'Iceland', 'Sweden',  
                                'Singapore', 'Denmark'])
```

ndarray for countries

```
In [18]: s_country = pd.Series(np_country)
```

Pass ndarray as an argument

```
In [19]: print (s_country)
```

```
0      Luxembourg
1         Norway
2          Japan
3    Switzerland
4   United States
5          Qatar
6         Iceland
7          Sweden
8     Singapore
9         Denmark
dtype: object
```

countries

Data type

Create Series from Scalar

In [31]: *#Print Series with scalar input*
`scalar_series = pd.Series(5., index=['a', 'b', 'c', 'd', 'e'])`

Scalar input

In [32]: `scalar_series`

Index

Out[32]:

a	5
b	5
c	5
d	5
e	5
dtype: float64	

index

Data

Data type

Accessing Elements in Series

Data can be accessed through different functions like loc, iloc by passing data element position or index range.

```
In [43]: #access elements in the series  
dict_country_gdp[0]
```

```
Out[43]: 52056.017809999998
```

```
In [44]: #access first 5 countries from the series  
dict_country_gdp[0:5]
```

```
Out[44]: Luxembourg      52056.01781  
Macao, China      40258.80862  
Norway      40034.85063  
Japan      39578.07441  
Switzerland      39170.41371  
dtype: float64
```

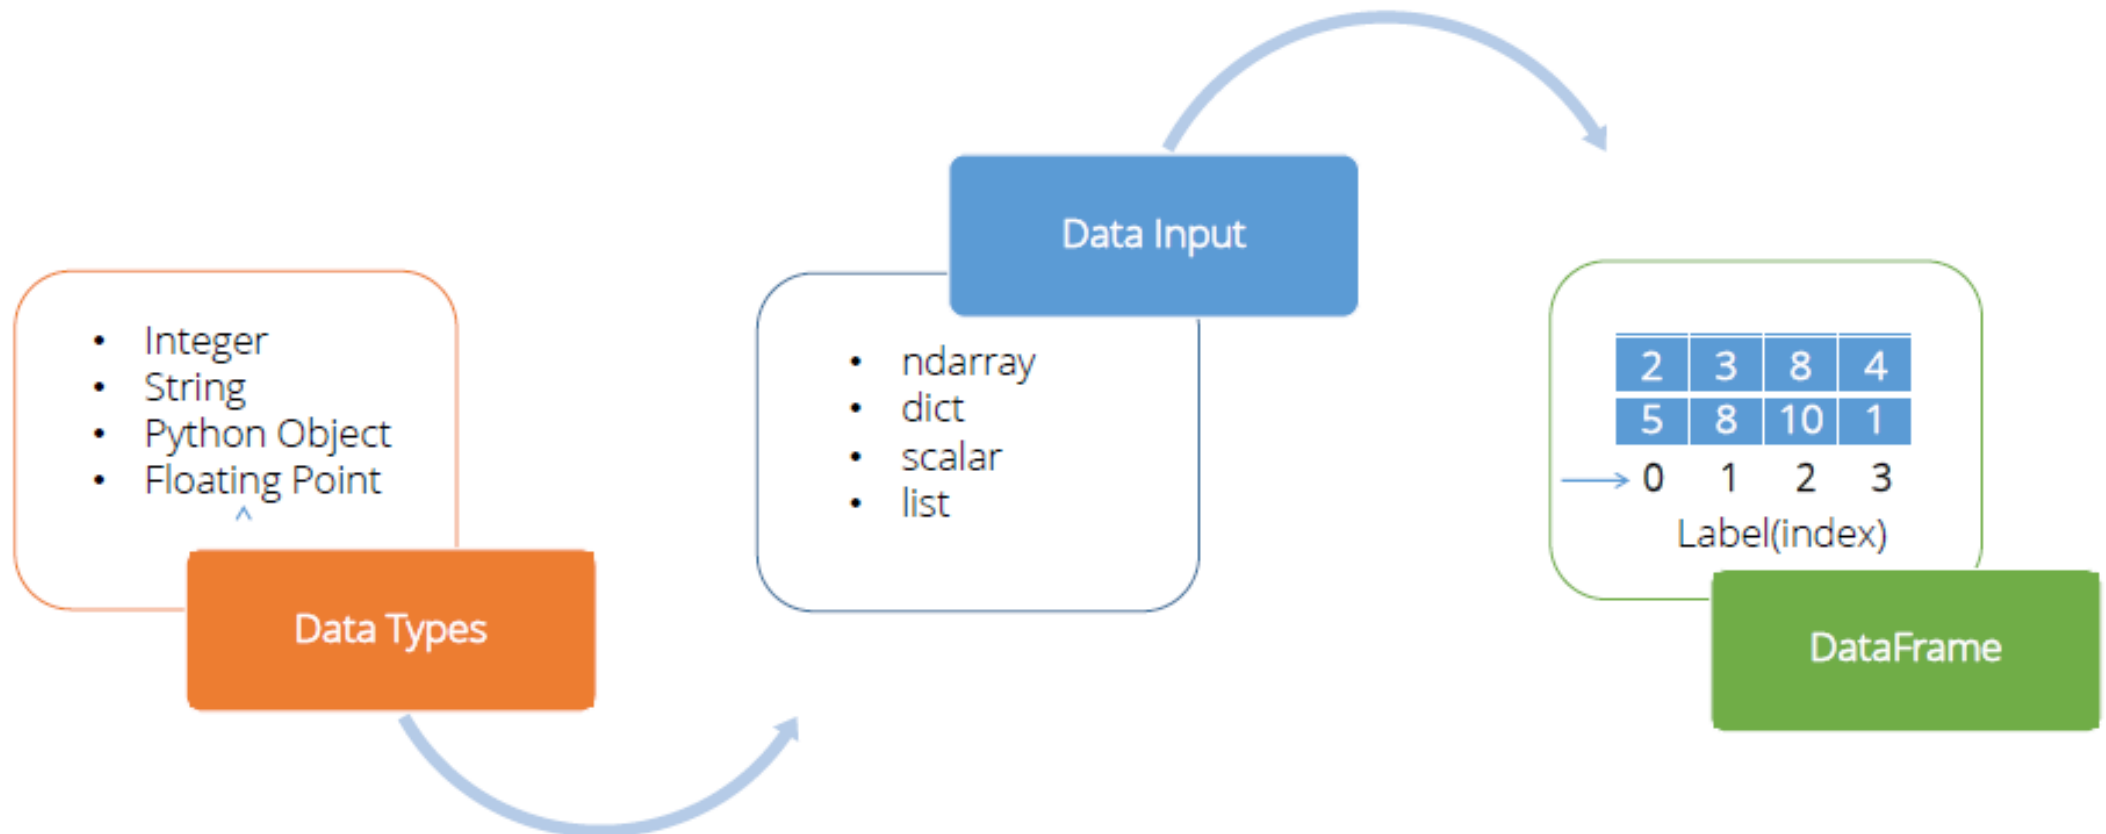
```
In [45]: #Look up a country by name or index  
dict_country_gdp.loc['United States']
```

```
Out[45]: 37691.027329999997
```

```
In [46]: #look up by position  
dict_country_gdp.iloc[0]
```

```
Out[46]: 52056.017809999998
```

DataFrame is a two-dimensional labeled data structure with columns of potentially different types.



Create DataFrame from Lists

Let's see how you can create a DataFrame from lists:

```
In [1]: import pandas as pd
```

Create DataFrame from dict of equal length lists

```
In [2]: #Last five olympics data: place, year and number of countries participated  
olympic_data_list = {'HostCity': ['London', 'Beijing', 'Athens', 'Sydney', 'Atlanta'],  
                     'Year': [2012, 2008, 2004, 2000, 1996],  
                     'No. of Participating Countries': [205, 204, 201, 200, 197]  
                     }
```

```
In [3]: df_olympic_data = pd.DataFrame(olympic_data_list) ← Pass the list to the DataFrame
```

```
In [4]: df_olympic_data
```

```
Out[4]:
```

	HostCity	No. of Participating Countries	Year
0	London	205	2012
1	Beijing	204	2008
2	Athens	201	2004
3	Sydney	200	2000
4	Atlanta	197	1996

Create DataFrame from dict

This example shows you how to create a DataFrame from a series of dicts:

Create DataFrame from dict of dicts

In [5]: `olympic_data_dict = {'London':{2012:205}, 'Beijing':{2008:204}}`



In [6]: `df_olympic_data_dict = pd.DataFrame(olympic_data_dict)`

In [7]: `df_olympic_data_dict`

Out[7]:

	Beijing	London
2008	204	NaN
2012	NaN	205

View Data Frame

You can view a DataFrame by referring the column name or with the describe function.

```
In [8]: #select by City name  
df_olympic_data.HostCity
```

```
Out[8]: 0    London  
1    Beijing  
2    Athens  
3    Sydney  
4    Atlanta  
Name: HostCity, dtype: object
```

```
In [9]: #use describe function to display the content  
df_olympic_data.describe
```

```
Out[9]: <bound method DataFrame.describe of      HostCity  No. of Participating Countries  Year  
0    London      205      2012  
1   Beijing      204      2008  
2    Athens      201      2004  
3   Sydney      200      2000  
4   Atlanta      197      1996>
```

Create DataFrame from ndarray

Create DataFrame from dict of ndarray

```
In [13]: import numpy as np
```

```
In [14]: np_array = np.array([2012, 2008, 2004, 2006]) ← Create an ndarray with years  
dict_ndarray = {'year': np_array} ← Create a dict with the ndarray
```

```
In [15]: df_ndarray = pd.DataFrame(dict_ndarray) ← Pass this dict to a new DataFrame
```

```
In [16]: df_ndarray
```

```
Out[16]:
```

	year
0	2012
1	2008
2	2004
3	2006

Create DataFrame from DataFrame

Create DataFrame from DataFrame object

```
In [17]: df_from_df = pd.DataFrame(df_olympic_series)
```


Create a DataFrame from a DataFrame

```
In [18]: df_from_df
```

Out[18]:

	Host Cities	No. of Participating Countries
2012	London	205
2008	Beijing	204
2004	Athens	201
2000	Sydney	200
1996	Atlanta	197

How do I read and write tabular data?



?

```
In [2]: titanic = pd.read_csv("data/titanic.csv")
```

pandas provides the `read_csv()` function to read data stored as a csv file into a pandas `DataFrame`. pandas supports many different file formats or data sources out of the box (csv, excel, sql, json, parquet, ...), each of them with the prefix `read_*`.

How do I read and write tabular data?

Make sure to always have a check on the data after reading in the data. When displaying a `DataFrame`, the first and last 5 rows will be shown by default:

```
In [3]: titanic
Out[3]:
```

	PassengerId	Survived	Pclass	Name	Sex	...
0	1	0	3	Braund, Mr. Owen Harris	male	...
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	...
2	3	1	3	Heikkinen, Miss. Laina	female	...
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	...
4	5	0	3	Allen, Mr. William Henry	male	...
..
886	887	0	2	Montvila, Rev. Juozas	male	...
887	888	1	1	Graham, Miss. Margaret Edith	female	...
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	...
889	890	1	1	Behr, Mr. Karl Howell	male	...
890	891	0	3	Dooley, Mr. Patrick	male	...

[891 rows x 12 columns]

head()

```
In [4]: titanic.head(8)
```

```
Out[4]:
```

	PassengerId	Survived	Pclass	Name	Sex	...
0	1	0	3	Braund, Mr. Owen Harris	male	...
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	...
2	3	1	3	Heikkinen, Miss. Laina	female	...
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	...
4	5	0	3	Allen, Mr. William Henry	male	...
5	6	0	3	Moran, Mr. James	male	...
6	7	0	1	McCarthy, Mr. Timothy J	male	...
7	8	0	3	Palsson, Master. Gosta Leonard	male	...

```
[8 rows x 12 columns]
```

To see the first N rows of a `DataFrame`, use the `head()` method with the required number of rows (in this case 8) as argument.

Head() , tail()

Note

Interested in the last N rows instead? pandas also provides a `tail()` method. For example, `titanic.tail(10)` will return the last 10 rows of the DataFrame.

A check on how pandas interpreted each of the column data types can be done by requesting the pandas `dtypes` attribute:

```
In [5]: titanic.dtypes
Out[5]:
PassengerId      int64
Survived          int64
Pclass           int64
Name             object
Sex              object
Age             float64
SibSp            int64
Parch            int64
Ticket           object
Fare            float64
Cabin            object
Embarked         object
dtype: object
```

For each of the columns, the used data type is enlisted. The data types in this `DataFrame` are integers (`int64`), floats (`float64`) and strings (`object`).

Note

When asking for the `dtypes`, no brackets are used! `dtypes` is an attribute of a `DataFrame` and `Series`. Attributes of `DataFrame` or `Series` do not need brackets. Attributes represent a characteristic of a `DataFrame/Series`, whereas a method (which requires brackets) *do* something with the `DataFrame/Series` as introduced in the [first tutorial](#).

To_excel()

```
In [6]: titanic.to_excel('titanic.xlsx', sheet_name='passengers', index=False)
```

Whereas `read_*` functions are used to read data to pandas, the `to_*` methods are used to store data. The `to_excel()` method stores the data as an excel file. In the example here, the `sheet_name` is named *passengers* instead of the default *Sheet1*. By setting `index=False` the row index labels are not saved in the spreadsheet.

The equivalent read function `read_excel()` will reload the data to a `DataFrame`:

The method `info()` provides technical information about a `DataFrame`, so let's explain the output in more detail:

- It is indeed a `DataFrame`.
- There are 891 entries, i.e. 891 rows.
- Each row has a row label (aka the `index`) with values ranging from 0 to 890.
- The table has 12 columns. Most columns have a value for each of the rows (all 891 values are `non-null`). Some columns do have missing values and less than 891 `non-null` values.
- The columns `Name`, `Sex`, `Cabin` and `Embarked` consists of textual data (strings, aka `object`). The other columns are numerical data with some of them whole numbers (aka `integer`) and others are real numbers (aka `float`).
- The kind of data (characters, integers,...) in the different columns are summarized by listing the `dtypes`.
- The approximate amount of RAM used to hold the `DataFrame` is provided as well.

Remember

- Getting data in to pandas from many different file formats or data sources is supported by `read_*` functions.
- Exporting data out of pandas is provided by different `to_*` methods.
- The `head/tail/info` methods and the `dtypes` attribute are convenient for a first check.

How to select a subset of a DataFrame?

```
In [4]: ages = titanic["Age"]
```

```
In [5]: ages.head()
```

```
Out[5]:
```

```
0    22.0
```

```
1    38.0
```

```
2    26.0
```

```
3    35.0
```

```
4    35.0
```

```
Name: Age, dtype: float64
```

To select a single column, use square brackets `[]` with the column name of the column of interest.

Each column in a `DataFrame` is a `Series`. As a single column is selected, the returned object is a pandas `Series`.

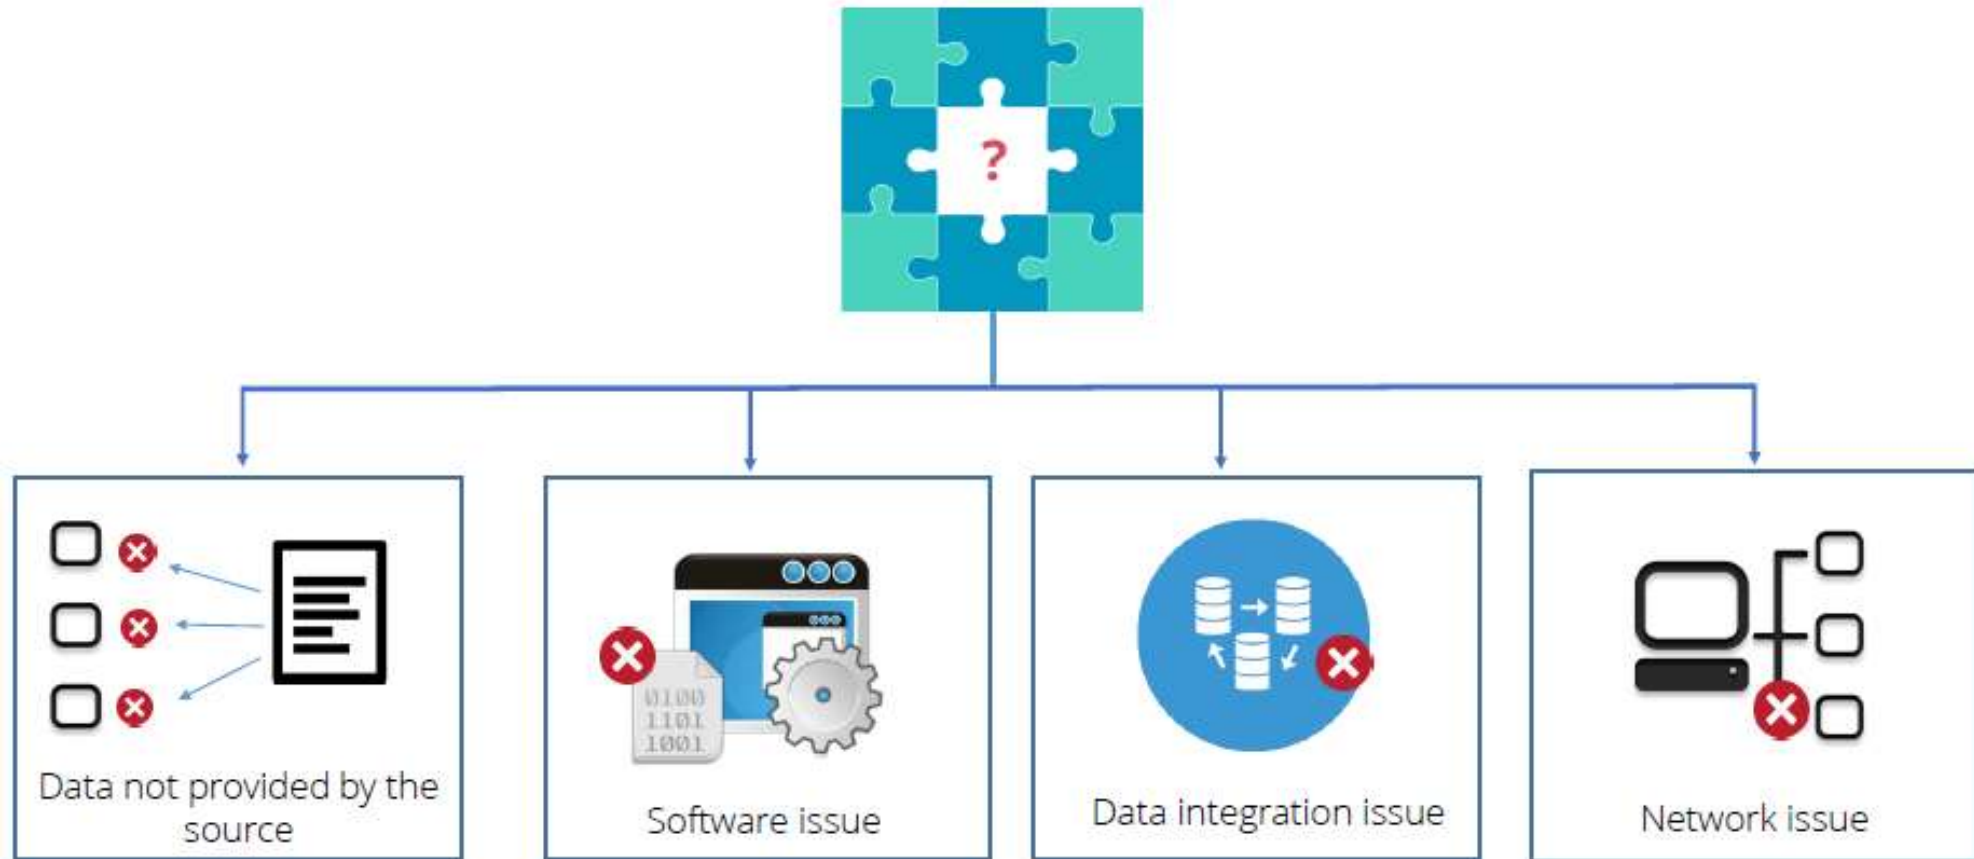
We can verify this by checking the type of the output:

```
In [6]: type(titanic["Age"])
```

```
Out[6]: pandas.core.series.Series
```

Missing Values

Various factors may lead to missing data values:



Handling Missing Values

It's difficult to operate on a dataset when it has missing values or uncommon indices.

```
In [3]: import pandas as pd
```

```
In [4]: #declare first series  
first_series = pd.Series([1,2,3,4,5],index=['a','b','c','d','e'])
```

```
In [5]: #declare second series  
second_series=pd.Series([10,20,30,40,50],index=['c','e','f','g','h'])
```

```
In [6]: sum_of_series = first_series+second_series
```

```
In [7]: sum_of_series
```

```
Out[7]: a    NaN  
       b    NaN  
       c    13  
       d    NaN  
       e    25  
       f    NaN  
       g    NaN  
       h    NaN  
       dtype: float64
```

Handling Missing Values with functions

The dropna function drops all the values with uncommon indices.

```
In [5]: sum_of_series
```

```
Out[5]: a      NaN  
       b      NaN  
       c     13.0  
       d      NaN  
       e     25.0  
       f      NaN  
       g      NaN  
       h      NaN  
       dtype: float64
```

```
In [6]: # drop NaN( Not a Number) values from dataset  
       dropna_s = sum_of_series.dropna()
```

```
In [7]: dropna_s
```

```
Out[7]: c     13.0  
       e     25.0  
       dtype: float64
```

Handling Missing Values with functions

The fillna function fills all the uncommon indices with a number instead of dropping them.

```
In [8]: dropna_s.fillna(0) ← Fill the missing values with zero
```

```
Out[8]: c    13.0  
        e    25.0  
        dtype: float64
```

```
In [9]: # Fill NaN( Not a Number) values with Zeroes (0)  
        fillna_s = sum_of_series.fillna(0) ←
```

```
In [10]: fillna_s
```

```
Out[10]: a    0.0  
         b    0.0  
         c    13.0  
         d    0.0  
         e    25.0  
         f    0.0  
         g    0.0  
         h    0.0  
         dtype: float64
```

Handling Missing Values with functions - example

```
In [10]: #fill values with zeroes before performing addition operation for missing indices  
fill_NaN_with_zeros_before_sum = first_series.add(second_series, fill_value=0)
```

```
In [11]: fill_NaN_with_zeros_before_sum
```

```
Out[11]: a      1  
b      2  
c     13  
d      4  
e     25  
f     30  
g     40  
h     50  
dtype: float64
```

Data operation can be performed through various built-in methods for faster data processing.

```
In [1]: import pandas as pd
```

```
In [2]: #declare movie rating dataframe: ratings from 1 to 5 (star * rating)
df_movie_rating = pd.DataFrame(
    {'movie 1': [5,4,3,3,2,1],
     'movie 2': [4,5,2,3,4,2]},
    index=['Tom', 'Jeff', 'Peter', 'Ram', 'Ted', 'Paul']
)
```

```
In [3]: df_movie_rating
```

```
Out[3]:
```

	movie 1	movie 2
Tom	5	4
Jeff	4	5
Peter	3	2
Ram	3	3
Ted	2	4
Paul	1	2

Data Operation with Statistical Functions

This example shows data operations with different statistical functions.

```
In [7]: df_test_scores = pd.DataFrame(  
        {'Test1': [95,84,73,88,82,61],  
        'Test2': [74,85,82,73,77,79]},  
        index=['Jack','Lewis','Patrick','Rich','Kelly','Paula']  
    )
```

← Create a DataFrame with two test

```
In [8]: df_test_scores.max()
```

← Apply the max function to find the maximum score

```
Out[8]: Test1    95  
        Test2    85  
        dtype: int64
```

```
In [9]: df_test_scores.mean()
```

← Apply the mean function to find the average score

```
Out[9]: Test1    80.500000  
        Test2    78.333333  
        dtype: float64
```

```
In [10]: df_test_scores.std()
```

← Apply the std function to find the standard deviation for both the tests

```
Out[10]: Test1    11.979149  
         Test2     4.633213  
         dtype: float64
```


Data Operation Using Groupby

This example shows how to operate data using the groupby function.

```
In [16]: df_president_name = pd.DataFrame({'first':['George','Bill', 'Ronald','Jimmy','George'],  
                                          'last':['Bush','Clinton', 'Regan', 'Carter', 'Washington']})
```

```
In [17]: df_president_name
```

```
Out[17]:
```

	first	last
0	George	Bush
1	Bill	Clinton
2	Ronald	Regan
3	Jimmy	Carter
4	George	Washington

Create a DataFrame with first and last name as former presidents

```
In [18]: grouped = df_president_name.groupby('first') ← Group the DataFrame with the first name
```

```
In [19]: grp_data = grouped.get_group('George') ← Group the DataFrame with the first name  
grp_data
```

```
Out[19]:
```

	first	last
0	George	Bush
4	George	Washington

Data Operation - Sorting

This example shows how to sort data

In [20]: `df_president_name.sort_values('first')` ← Sort values by first name

Out[20]:

	first	last
1	Bill	Clinton
0	George	Bush
4	George	Washington
3	Jimmy	Carter
2	Ronald	Regan



Which of the followings is used to store Two-dimensional data?

- a. Series
- b. DataFrame
- c. Panel
- d. PanelND

Which method is used for label-location indexing by label?

- a. iat
- b. iloc
- c. loc
- d. std

While viewing a dataframe, head() method will ____.

- a. return only the first row
- b. return only headers or column name of the DataFrame
- c. return the first five rows of the DataFrame
- d. throw an exception as it expects parameter(number) in parenthesis

How is an index for data elements assigned while creating a Pandas series/ select all that apply.

- a. Created automatically
- b. Needs to be assigned
- c. Once created can not be changed or altered
- d. Index is not applicable as series is one-dimensional

What will the result be in vector addition if label is not found in a series?

- a. Marked as Zeros for missing labels
- b. Labels will be skipped
- c. Marked as NaN for missing labels
- d. Will throw an exception, index not found