

Demo

- Go to view → choose Matrix
- Choose sales table and drag (Order Quantity) to the values field and (Product Key) to the rows fields.
- Replace the (Product Keys) rows with (Product Name) from (AW_product_lookup) table.

The screenshot shows the Power BI Data View interface. On the left, there is a list of tables: AW-Sales, AW_Calendar_Lookup, AW_Customer_Lookup, AW_Product_Categories_Lookup, and AW_Product_Lookup. The AW-Sales table is currently selected. Its schema is displayed on the right, showing columns: CustomerKey, OrderDate, OrderLineNumber, OrderQuantity, ProductKey, StockDate, and TerritoryKey. The OrderQuantity column is highlighted as the value field for the matrix. The ProductKey column is highlighted as the row field for the matrix.

The screenshot shows the Power BI Data View interface. The AW_Product_Lookup table is selected. Its schema is displayed on the right, showing columns: DiscountPrice, ModelName, ProductColor, ProductCost, ProductDescription, ProductKey (highlighted in yellow), ProductName (highlighted in red), ProductPrice, and ProductType. The ProductName column is highlighted as the row field for the matrix.

Note: We notice that we are limited to the scope of the sales table as there is no relationship with any Other table.

ProductKey	OrderQuantity
234	299
235	180
236	191
237	411
238	152
239	408
240	421
245	285
252	189
253	138
254	179
255	142
Total	84174

ProductName	OrderQuantity
All-Purpose Bike Stand	84174
AWC Logo Cap	84174
Silk Warm - Downer	84174
Globe Lock	84174
Oven	84174
Classic Vest L	84174
Classic Vest M	84174
Classic Vest S	84174
Front Brakes	84174
Rear Derailleur	84174
Kid's Mountain Bike	84174
Total	84174

DATABASE NORMALIZATION

Normalization is the process of organizing the tables and columns in a relational database to reduce redundancy and preserve data integrity. It's commonly used to:

- Eliminate redundant data to decrease table sizes and improve processing speed & efficiency
- Minimize errors and anomalies from data modifications (inserting, updating or deleting records)
- Simplify queries and structure the database for meaningful analysis

TIP: In a normalized database, each table should serve a *distinct* and *specific* purpose (i.e. product information, dates, transaction records, customer attributes, etc.)

date	product_id	quantity	product_brand	product_name	product_sku	product_weight
1/1/1997	869	5	Nationeel	Nationeel Grape Fruit Roll	52382137179	17
1/7/1997	869	2	Nationeel	Nationeel Grape Fruit Roll	52382137179	17
1/3/1997	1	4	Washington	Washington Berry Juice	90748583674	8.39
1/1/1997	1472	3	Fort West	Fort West Fudge Cookies	37276054024	8.28
1/6/1997	1472	2	Fort West	Fort West Fudge Cookies	37276054024	8.28
1/5/1997	2	4	Washington	Washington Mango Drink	96516502499	7.42
1/1/1997	76	4	Red Spade	Red Spade Sliced Chicken	62054644227	18.1
1/1/1997	76	2	Red Spade	Red Spade Sliced Chicken	62054644227	18.1
1/5/1997	3	2	Washington	Washington Strawberry Drink	58427771925	13.1
1/7/1997	3	2	Washington	Washington Strawberry Drink	58427771925	13.1
1/1/1997	320	3	Excellent	Excellent Cranberry Juice	36570182442	16.4

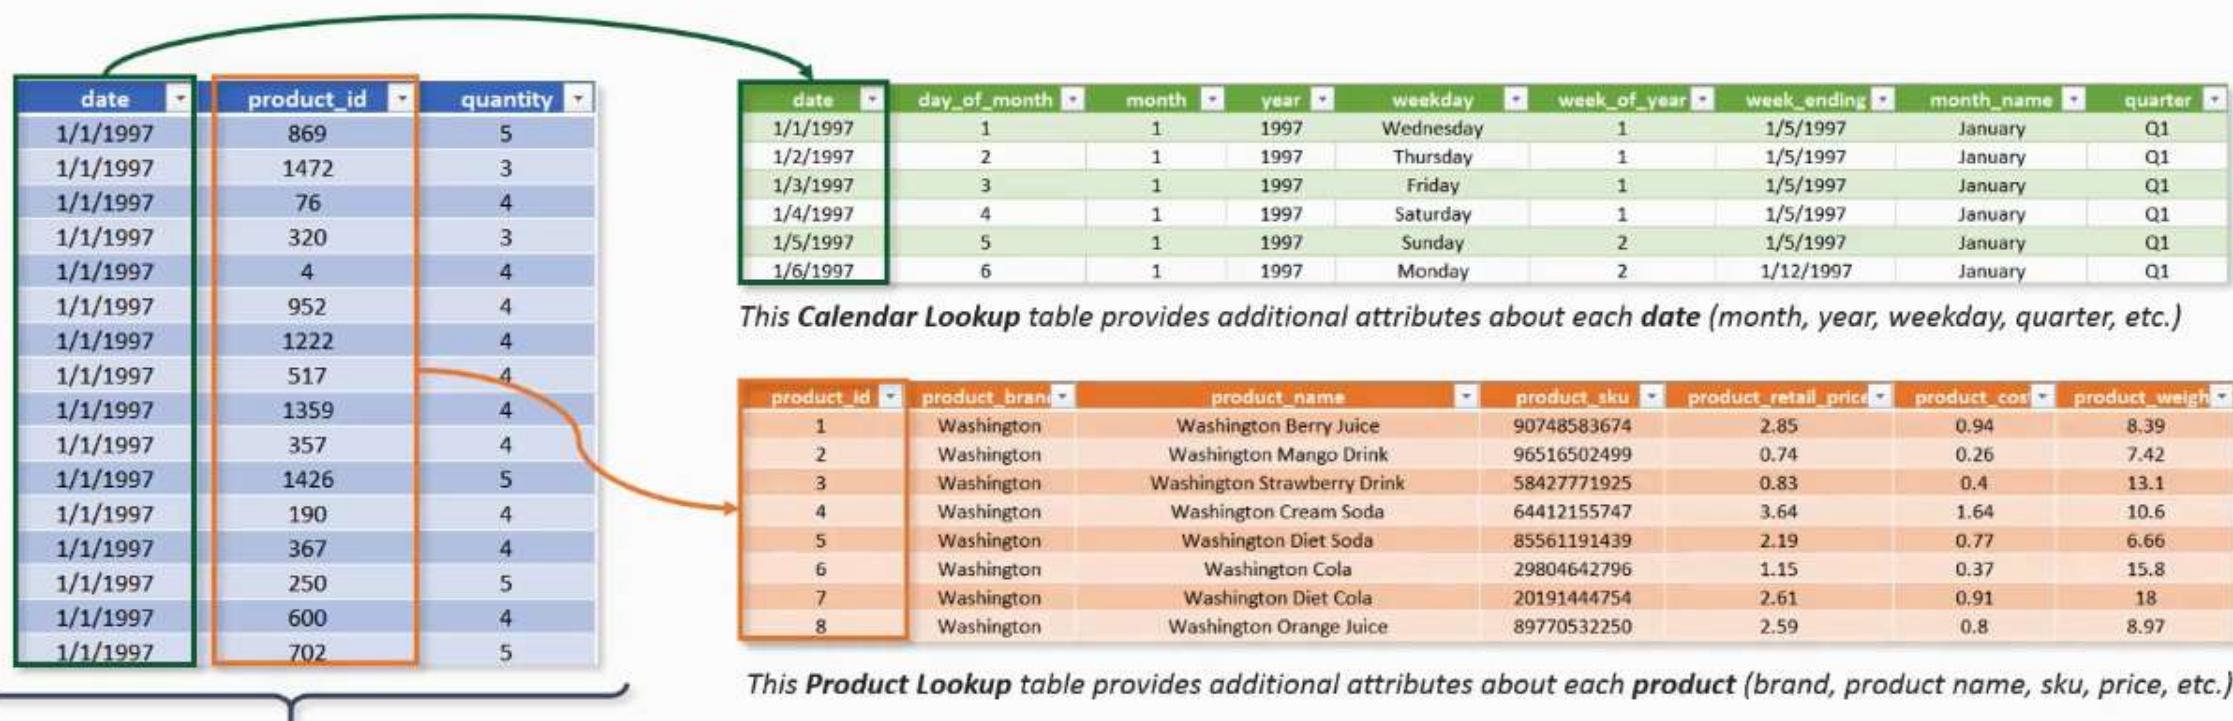
When you **don't normalize**, you end up with tables like this; all of the rows with duplicate product info could be eliminated with a lookup table based on **product_id**

This may not seem critical now, but minor inefficiencies can become major problems as databases scale in size!

DATA TABLES VS. LOOKUP TABLES

Models generally contain two types of tables: **data** (or “*fact*”) tables, and **lookup** (or “*dimension*”) tables

- **Data tables** contain *numbers* or *values*, typically at a granular level, with ID or “*key*” columns that can be used to create table relationships
- **Lookup tables** provide descriptive, often text-based *attributes* about each dimension in a table



This Data Table contains “quantity” values, and connects to lookup tables via the “date” and “product_id” columns

PRIMARY VS. FOREIGN KEYS

date	product_id	quantity
1/1/1997	869	5
1/1/1997	1472	3
1/1/1997	76	4
1/1/1997	320	3
1/1/1997	4	4
1/1/1997	952	4
1/1/1997	1222	4
1/1/1997	517	4
1/1/1997	1359	4
1/1/1997	357	4
1/1/1997	1426	5
1/1/1997	190	4
1/1/1997	367	4
1/1/1997	250	5
1/1/1997	600	4
1/1/1997	702	5

date	day_of_month	month	year	weekday	week_of_year	week_ending	month_name	quarter
1/1/1997	1	1	1997	Wednesday	1	1/5/1997	January	Q1
1/2/1997	2	1	1997	Thursday	1	1/5/1997	January	Q1
1/3/1997	3	1	1997	Friday	1	1/5/1997	January	Q1
1/4/1997	4	1	1997	Saturday	1	1/5/1997	January	Q1
1/5/1997	5	1	1997	Sunday	2	1/5/1997	January	Q1
1/6/1997	6	1	1997	Monday	2	1/12/1997	January	Q1

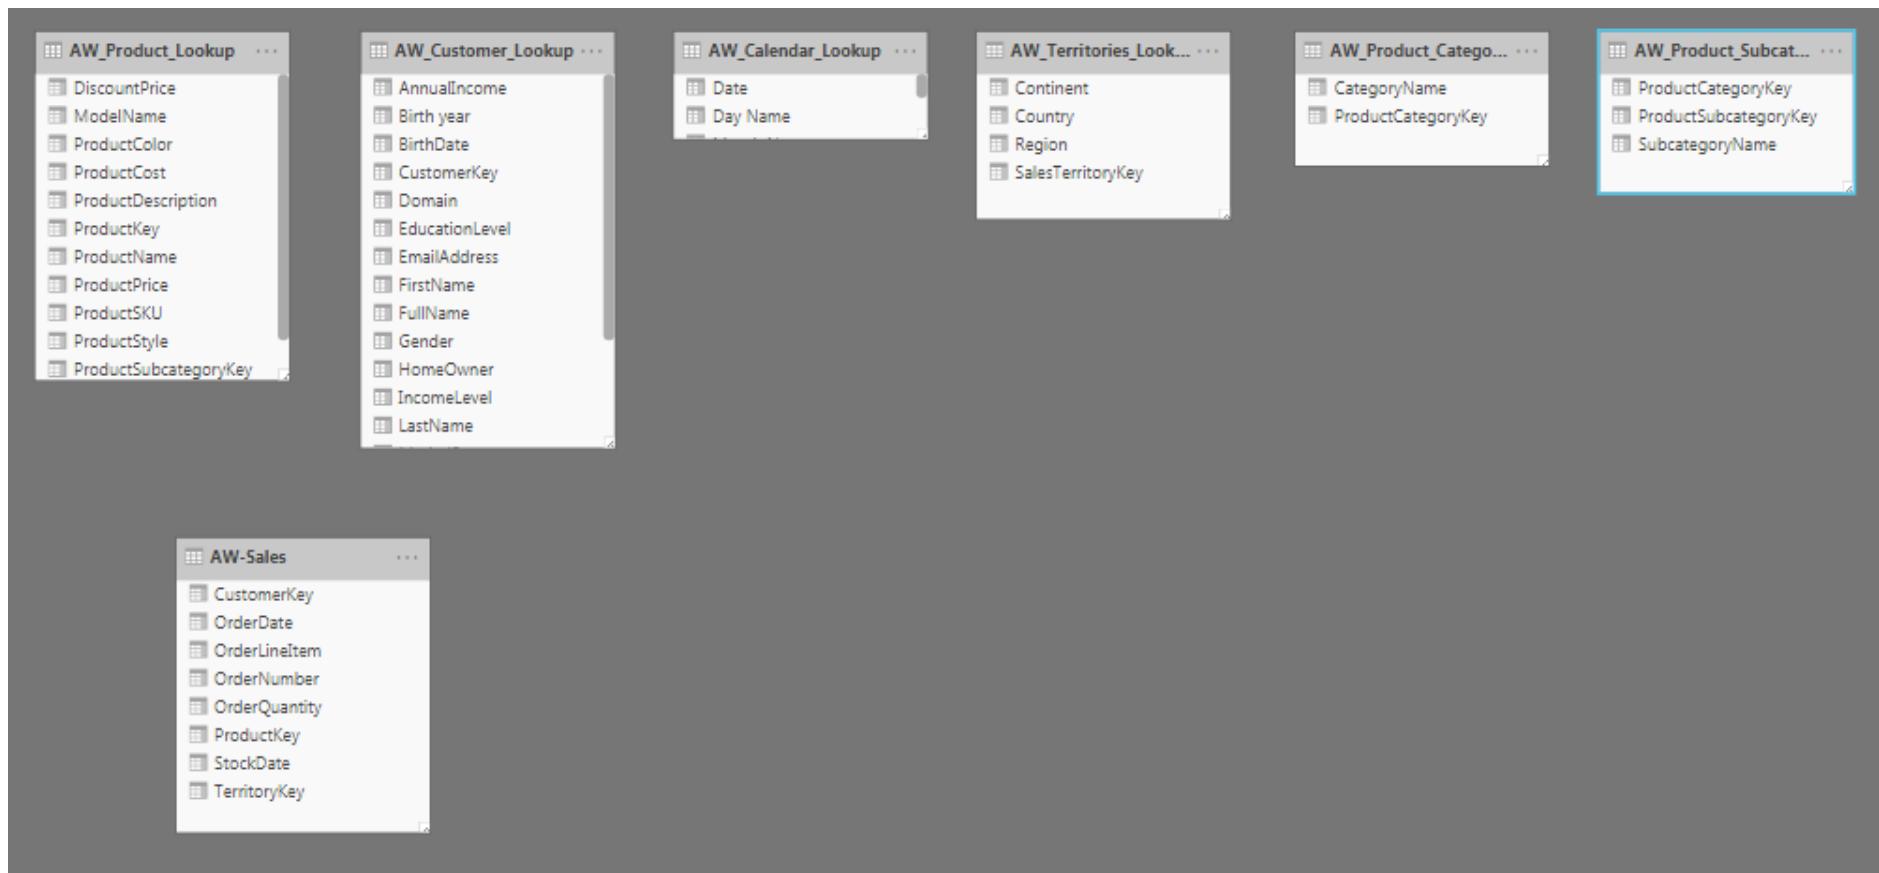
product_id	product_brand	product_name	product_sku	product_retail_price	product_cost	product_weight
1	Washington	Washington Berry Juice	90748583674	2.85	0.94	8.39
2	Washington	Washington Mango Drink	96516502499	0.74	0.26	7.42
3	Washington	Washington Strawberry Drink	58427771925	0.83	0.4	13.1
4	Washington	Washington Cream Soda	64412155747	3.64	1.64	10.6
5	Washington	Washington Diet Soda	85561191439	2.19	0.77	6.66
6	Washington	Washington Cola	29804642796	1.15	0.37	15.8
7	Washington	Washington Diet Cola	20191444754	2.61	0.91	18
8	Washington	Washington Orange Juice	89770532250	2.59	0.8	8.97

These columns are **foreign keys**; they contain *multiple* instances of each value, and are used to match the **primary keys** in related lookup tables

These columns are **primary keys**; they *uniquely* identify each row of a table, and match the **foreign keys** in related data tables

Demo

- Let's divide our tables in power BI model tab to lookup and data table.



RELATIONSHIPS VS. MERGED TABLES



*Can't I just **merge queries** or use **LOOKUP** or **RELATED** functions to pull those attributes into the fact table itself, so that I have everything in one place??*

-Anonymous confused man

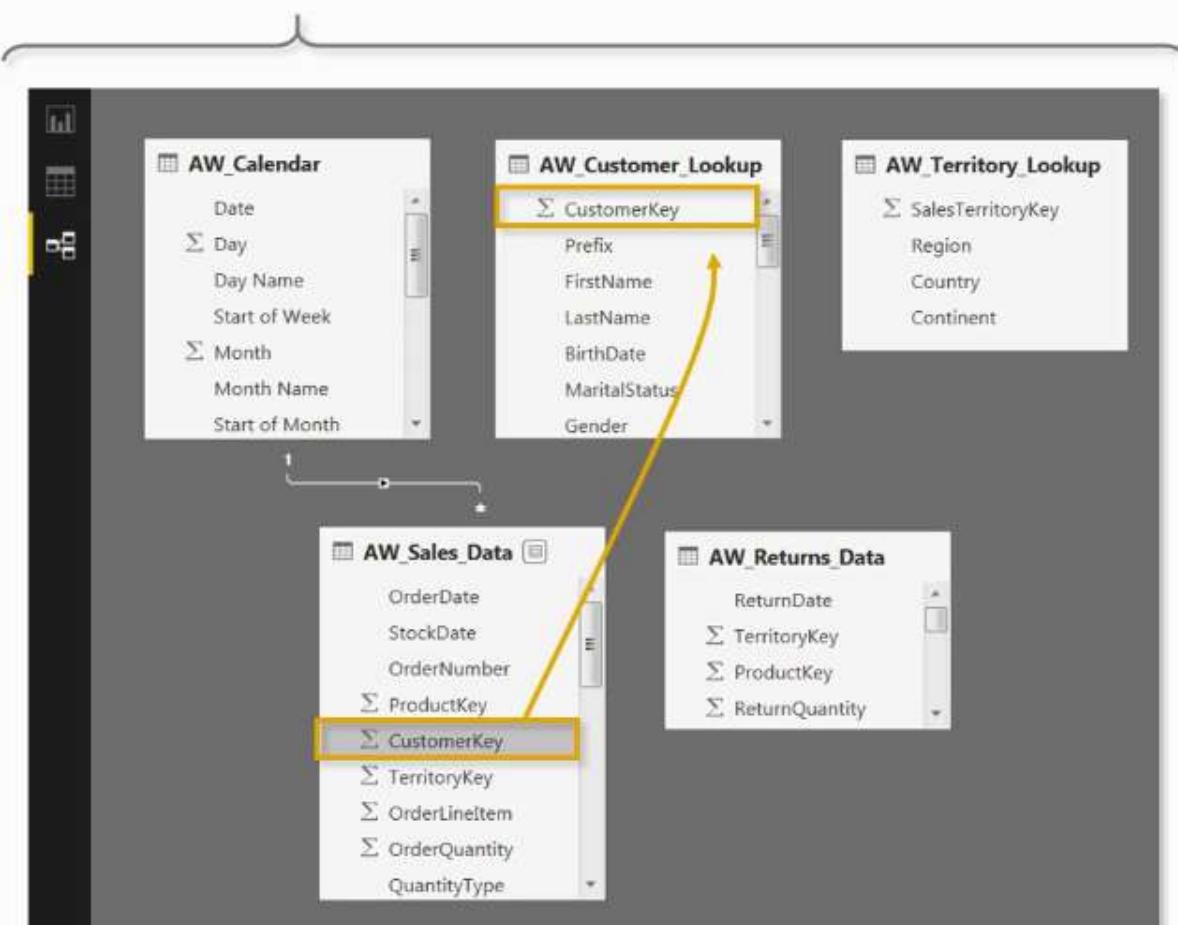
Original Fact Table fields			Attributes from Calendar Lookup table								Attributes from Product Lookup table			
date	product_id	quantity	day_of_month	month	year	weekday	month_name	quarter	product_brand	product_name	product_sku	product_weight		
1/1/1997	869	5	1	1	1997	Wednesday	January	Q1	Nationeel	Nationeel Grape Fruit Roll	52382137179	17		
1/7/1997	869	2	7	1	1997	Tuesday	January	Q1	Nationeel	Nationeel Grape Fruit Roll	52382137179	17		
1/3/1997	1	4	3	1	1997	Friday	January	Q1	Washington	Washington Berry Juice	90748583674	8.39		
1/1/1997	1472	3	1	1	1997	Wednesday	January	Q1	Fort West	Fort West Fudge Cookies	37276054024	8.28		
1/6/1997	1472	2	6	1	1997	Monday	January	Q1	Fort West	Fort West Fudge Cookies	37276054024	8.28		
1/5/1997	2	4	5	1	1997	Sunday	January	Q1	Washington	Washington Mango Drink	96516502499	7.42		
1/1/1997	76	4	1	1	1997	Wednesday	January	Q1	Red Spade	Red Spade Sliced Chicken	62054644227	18.1		
1/1/1997	76	2	1	1	1997	Wednesday	January	Q1	Red Spade	Red Spade Sliced Chicken	62054644227	18.1		
1/5/1997	3	2	5	1	1997	Sunday	January	Q1	Washington	Washington Strawberry Drink	58427771925	13.1		
1/7/1997	3	2	7	1	1997	Tuesday	January	Q1	Washington	Washington Strawberry Drink	58427771925	13.1		
1/1/1997	320	3	1	1	1997	Wednesday	January	Q1	Excellent	Excellent Cranberry Juice	36570182442	16.4		

Sure you can, **but it's inefficient!**

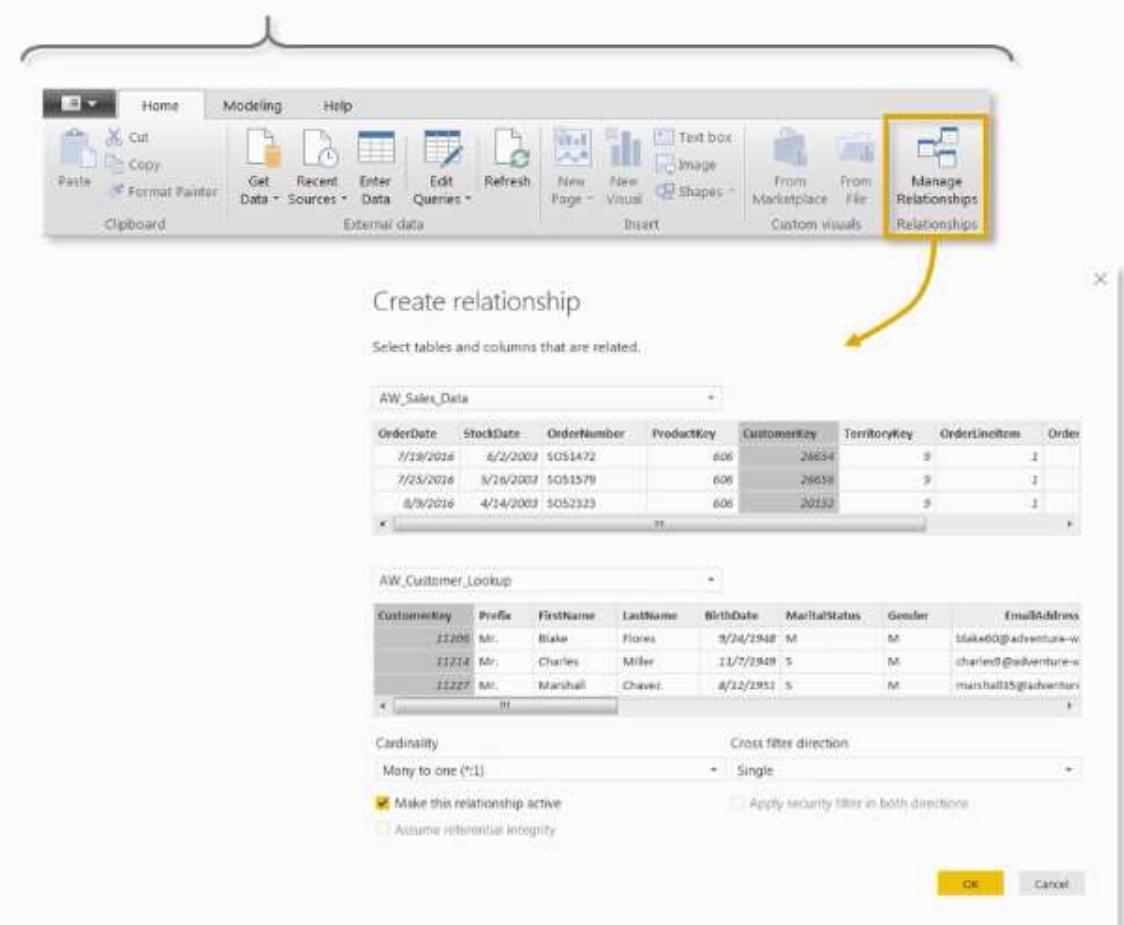
- Merging data in this way creates **redundant data** and utilizes **significantly more memory and processing power** than creating relationships between multiple small tables

CREATING TABLE RELATIONSHIPS

Option 1: Click and drag to connect primary and foreign keys within the **Relationships** pane



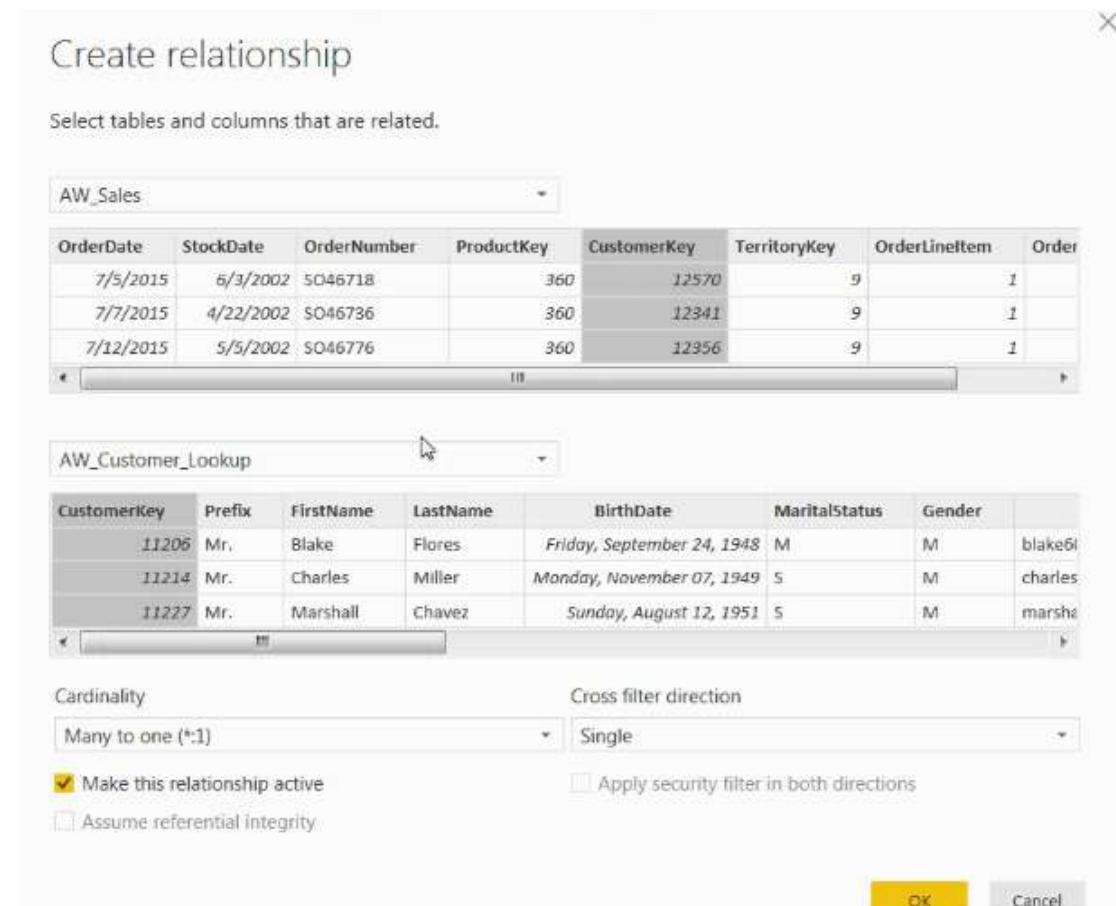
Option 2: Add or detect relationships using the “**Manage Relationships**” dialog box



Demo : Create a Relationship using Manage Relationships

- Go to model tab → click (Manage Relationship) from the ribbon.
- Click new
- Choose your columns to make a Relationship.

Note: The relationship was already Created as both columns names are same.



Demo : Create a Relationship by drag and drop

- Got to model tab
- Drag (orderdate) from sales table to date from (calender_lookup) to create a manual relationship.
- Drag (Territorieskey) from sales table to (SalesTerritoriesKey) from (Territories_lookup) to create a manual relationship.
- Drag (Productkey) from sales table to (Productkey) from (Product_lookup) to create a manual relationship.
- Now check from the view/report tab, you will find that The (Order quantity values) are divided based on (product name)



Join Data - Join Multiple Datasets based on common keys



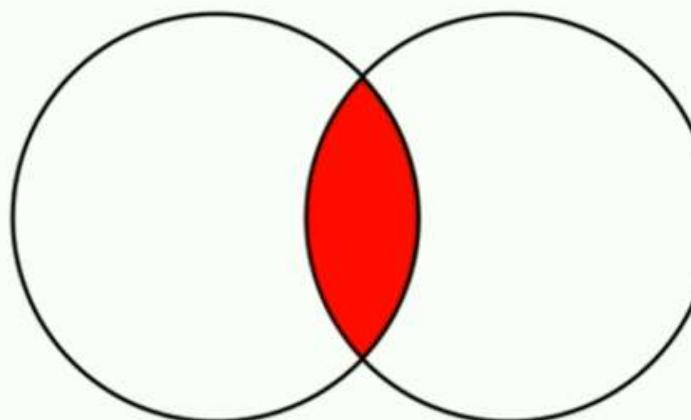
- Information is provided in two or more datasets
 - Different sources
 - Created at different times
- Datasets are related by key columns
- Different types of Join supported by AzureML
 - Inner Join
 - Left Outer Join
 - Full Outer Join
 - Left Semi-join

Inner Join



EmpID	Salary
EMP001	\$ 5,000
EMP002	\$ 5,500
EMP003	\$ 5,200
EMP004	\$ 6,000
EMP007	\$ 5,800
EMP008	\$ 6,700

EmpID	Department
EMP001	IT
EMP003	IT
EMP004	Marketing
EMP007	Finance
EMP009	Marketing
EMP010	Finance



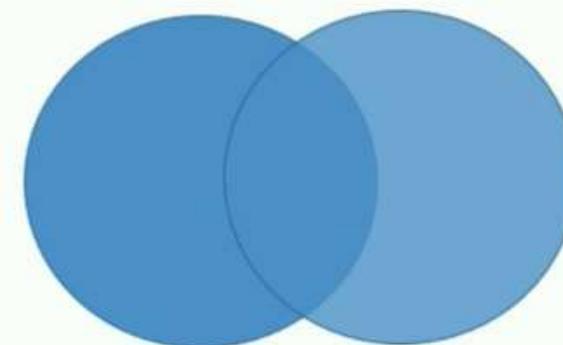
EmpID	Salary	Department
EMP001	\$ 5,000	IT
EMP003	\$ 5,200	IT
EMP004	\$ 6,000	Marketing
EMP007	\$ 5,800	Finance

Full Outer Join



EmplID	Salary
EMP001	\$ 5,000
EMP002	\$ 5,500
EMP003	\$ 5,200
EMP004	\$ 6,000
EMP007	\$ 5,800
EMP008	\$ 6,700

EmplID	Department
EMP001	IT
EMP003	IT
EMP004	Marketing
EMP007	Finance
EMP009	Marketing
EMP010	Finance



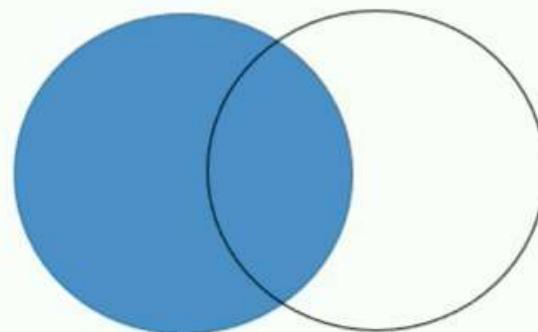
EmplID	Salary	Department
EMP001	\$ 5,000	IT
EMP002	\$ 5,500	
EMP003	\$ 5,200	IT
EMP004	\$ 6,000	Marketing
EMP007	\$ 5,800	Finance
EMP008	\$ 6,700	
EMP009		Marketing
EMP010		Finance

Left Outer Join



EmplID	Salary
EMP001	\$ 5,000
EMP002	\$ 5,500
EMP003	\$ 5,200
EMP004	\$ 6,000
EMP007	\$ 5,800
EMP008	\$ 6,700

EmplID	Department
EMP001	IT
EMP003	IT
EMP004	Marketing
EMP007	Finance
EMP009	Marketing
EMP010	Finance



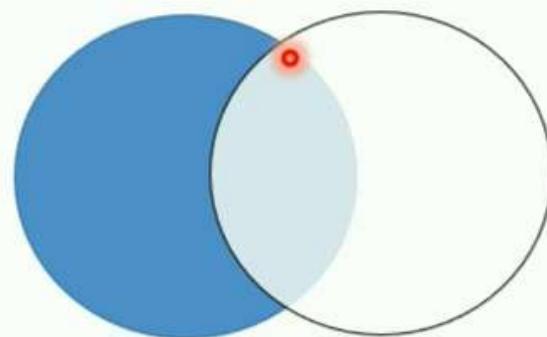
EmplID	Salary	Department
EMP001	\$ 5,000	IT
EMP002	\$ 5,500	
EMP003	\$ 5,200	IT
EMP004	\$ 6,000	Marketing
EMP007	\$ 5,800	Finance
EMP008	\$ 6,700	

Left Semi Join

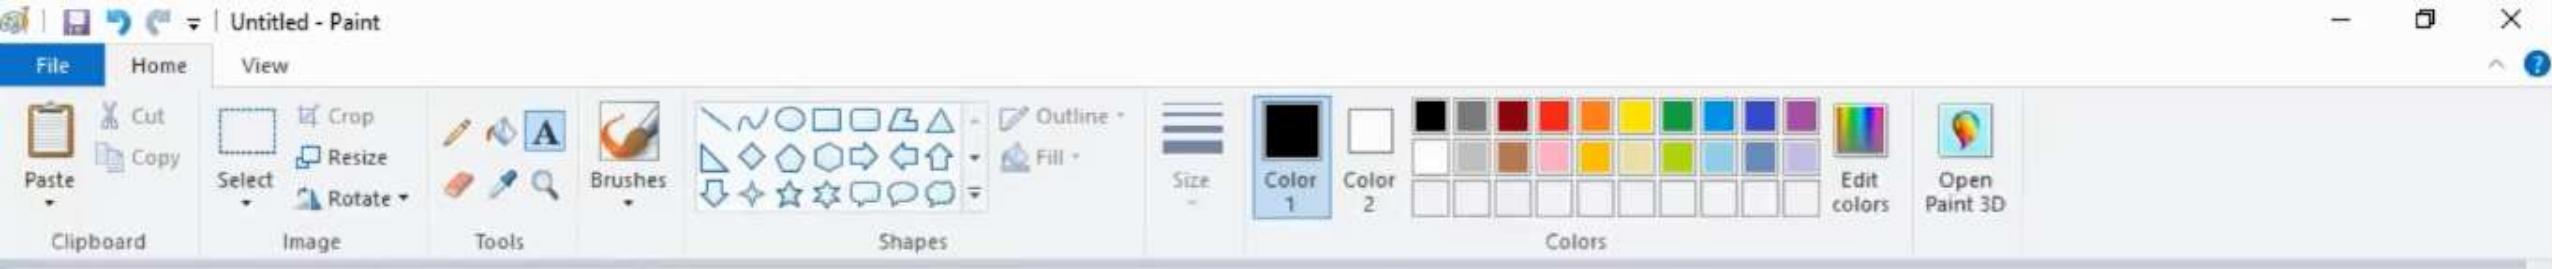


EmpID	Salary
EMP001	\$ 5,000
EMP002	\$ 5,500
EMP003	\$ 5,200
EMP004	\$ 6,000
EMP007	\$ 5,800
EMP008	\$ 6,700

EmpID	Department
EMP001	IT
EMP003	IT
EMP004	Marketing
EMP007	Finance
EMP009	Marketing
EMP010	Finance

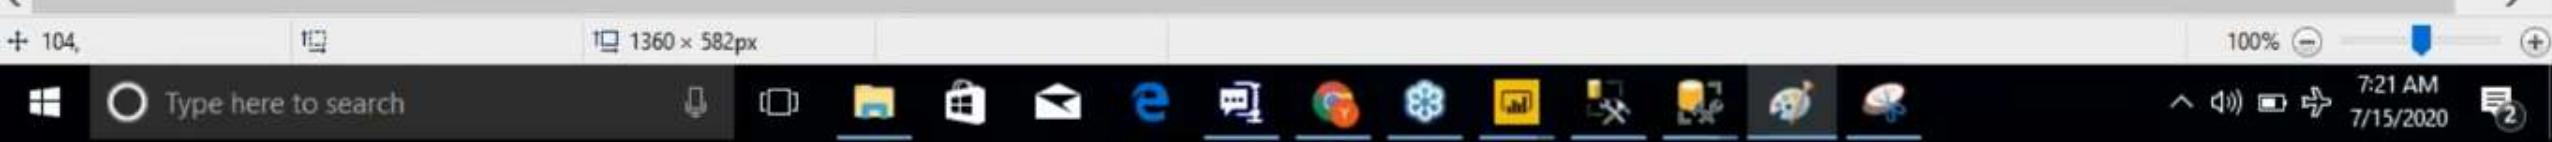


EmpID	Salary
EMP001	\$ 5,000
EMP003	\$ 5,200
EMP004	\$ 6,000
EMP007	\$ 5,800

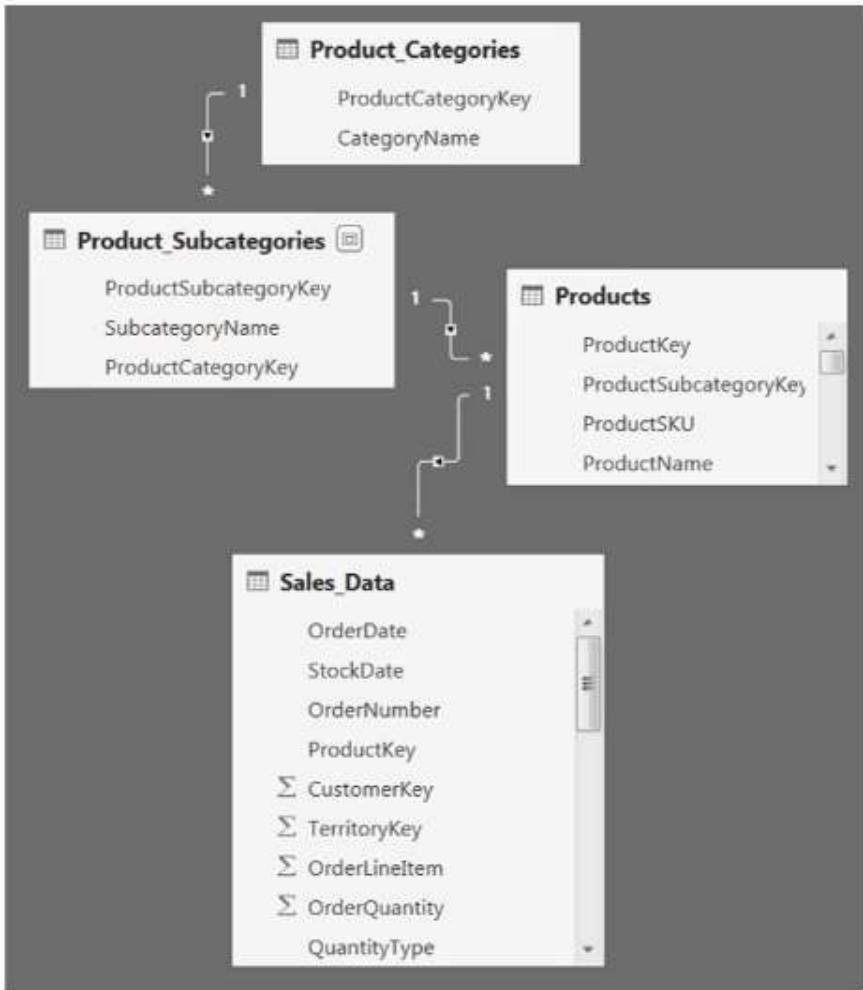


EID	empname	Salary	DateOfJoin	JoiningBonus	DID
00004001	Sone	500000	1/1/2020 12:00:00 AM	50000	1
0001001	Kumar	500000	1/1/2020 12:00:00 AM	50000	2
0002001	Tom	500000	1/1/2020 12:00:00 AM	50000	4
0003001	Hary	500000	1/1/2020 12:00:00 AM	50000	null
01920	Imran	200000	1/3/2020 12:00:00 AM	30000	2
E0101	Yousuf	100000	1/1/2020 12:00:00 AM	10000	1
E0102	Paula	100000	1/1/2020 12:00:00 AM	10000	null
E0103	Paula	100000	1/1/2020 12:00:00 AM	null	2
E0104	muhobath	100000	1/1/2020 12:00:00 AM	10000	1

DepID	DepName
1	IT
2	HR
3	Audit
null	General



CREATING “SNOWFLAKE” SCHEMAS



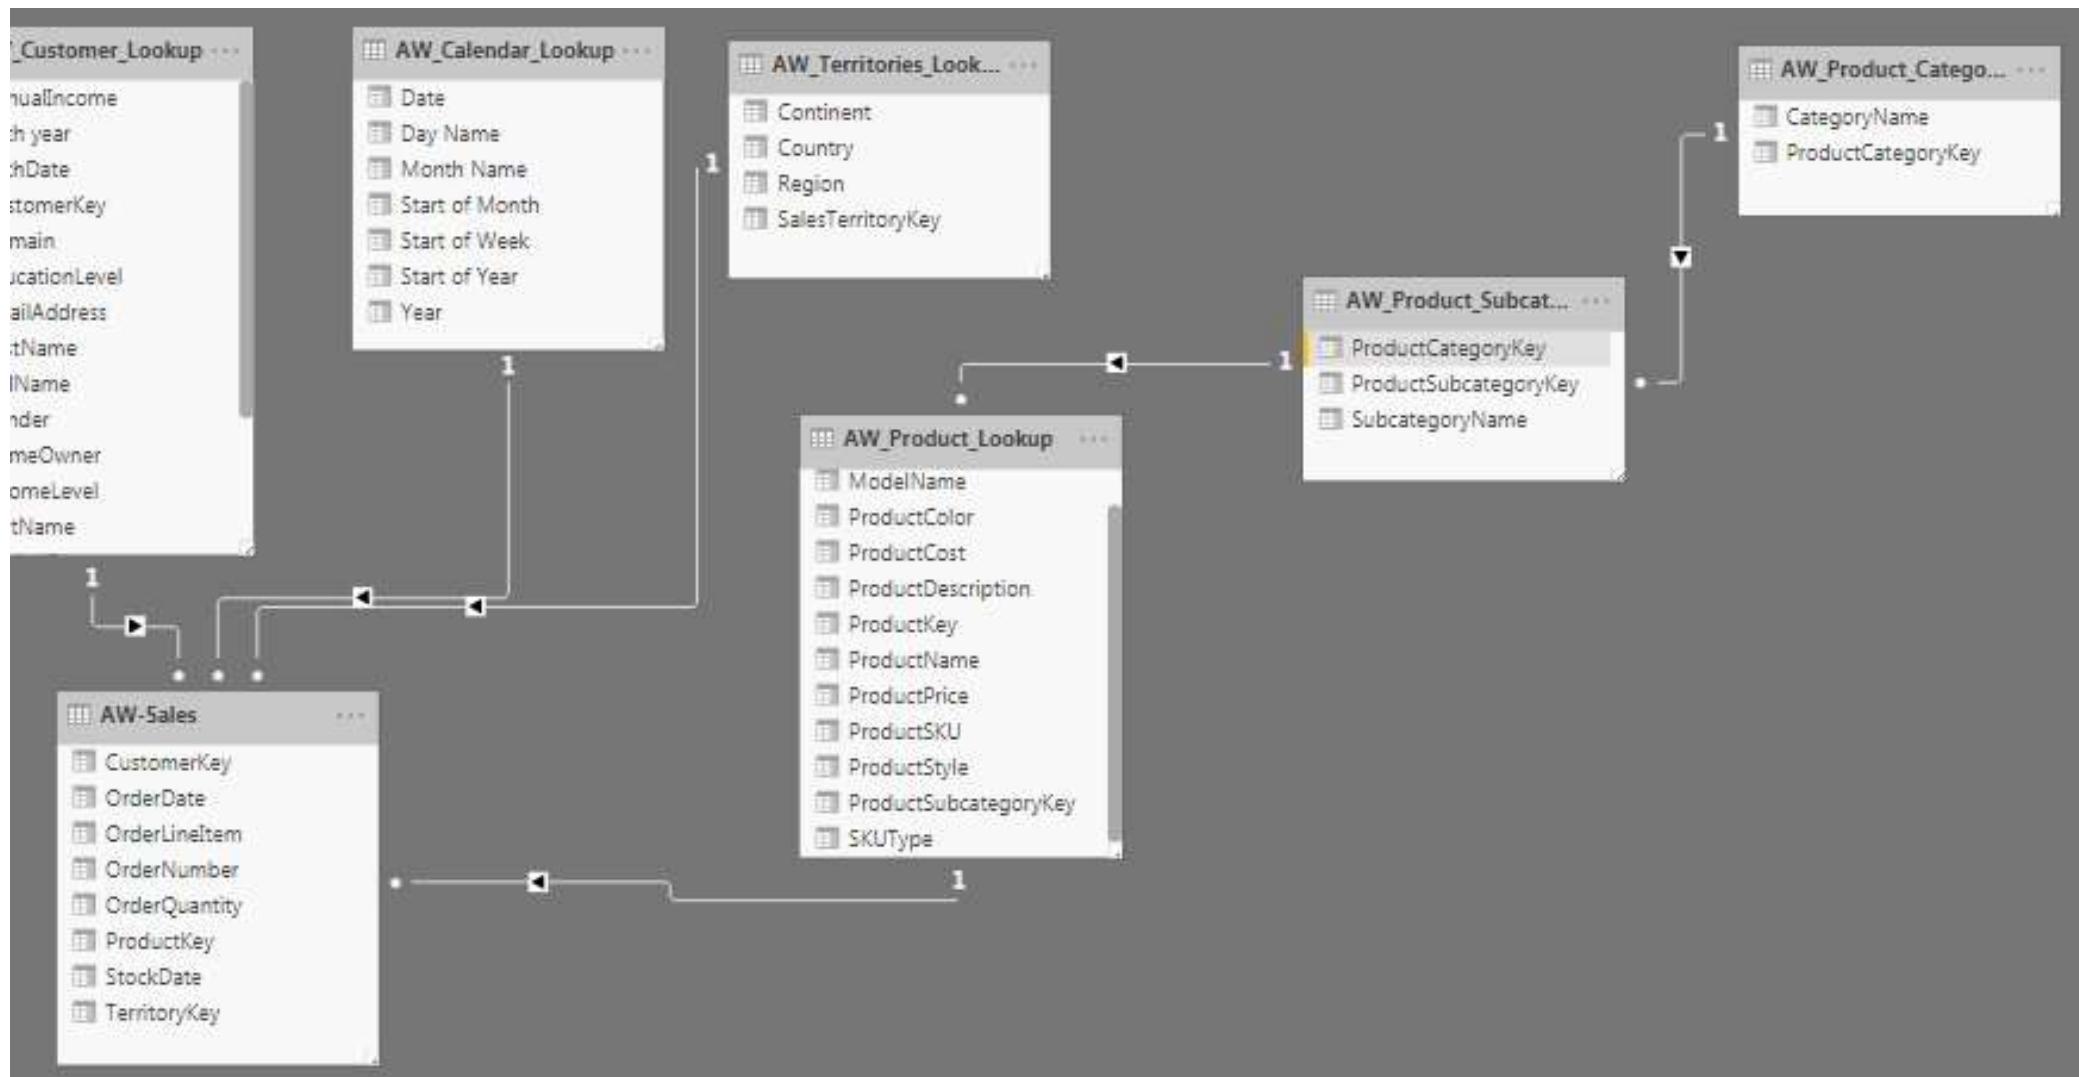
The **Sales_Data** table can connect to **Products** using the **ProductKey** field, but cannot connect directly to the **Subcategories** or **Categories** tables

By creating relationships from **Products** to **Subcategories** (using **ProductSubcategoryKey**) and **Subcategories** to **Categories** (using **ProductCategoryKey**), we have essentially connected **Sales_Data** to each lookup table; filter context will now flow all the way down the chain

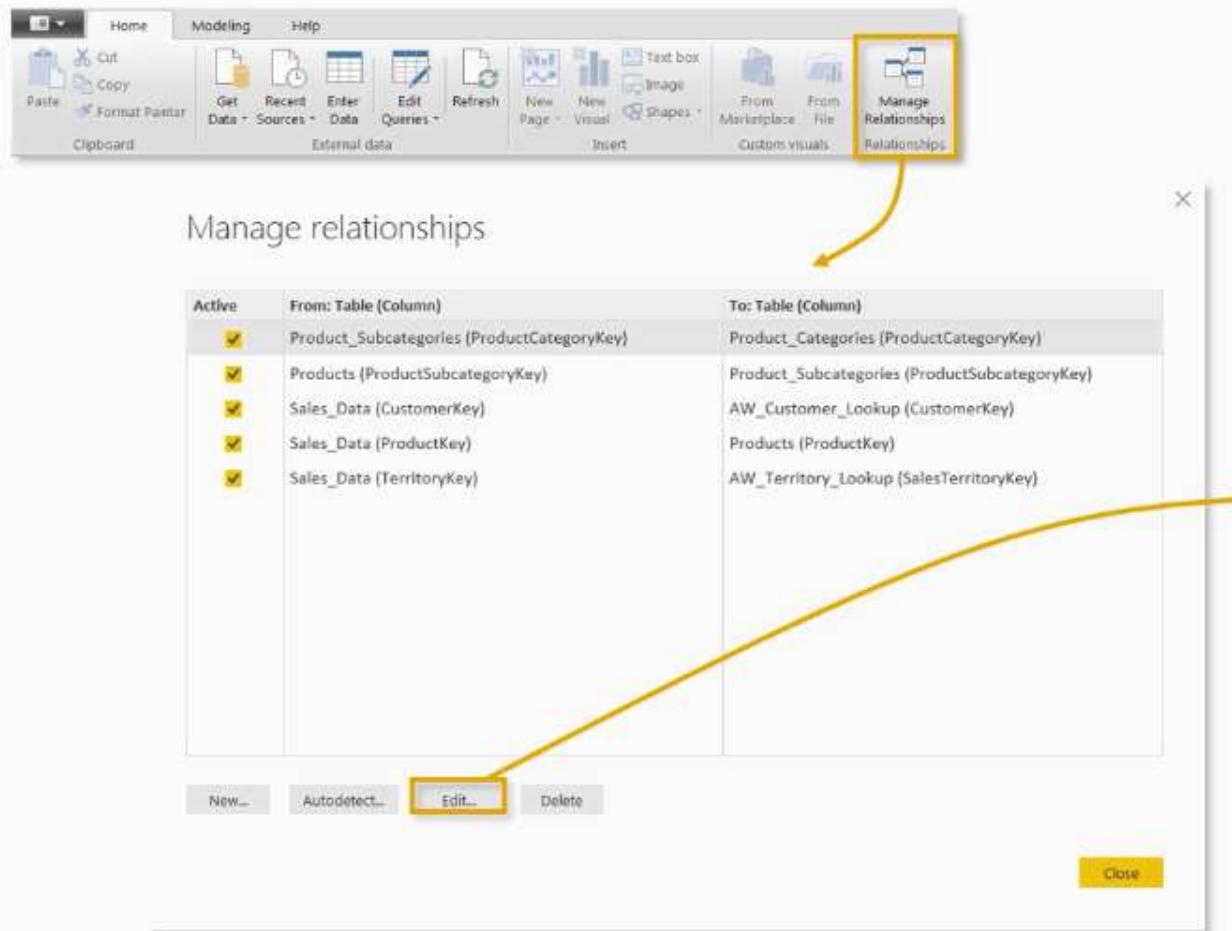
PRO TIP:

Models with chains of dimension tables are often called “snowflake” schemas (whereas “star” schemas usually have individual lookup tables surrounding a central data table)

Demo



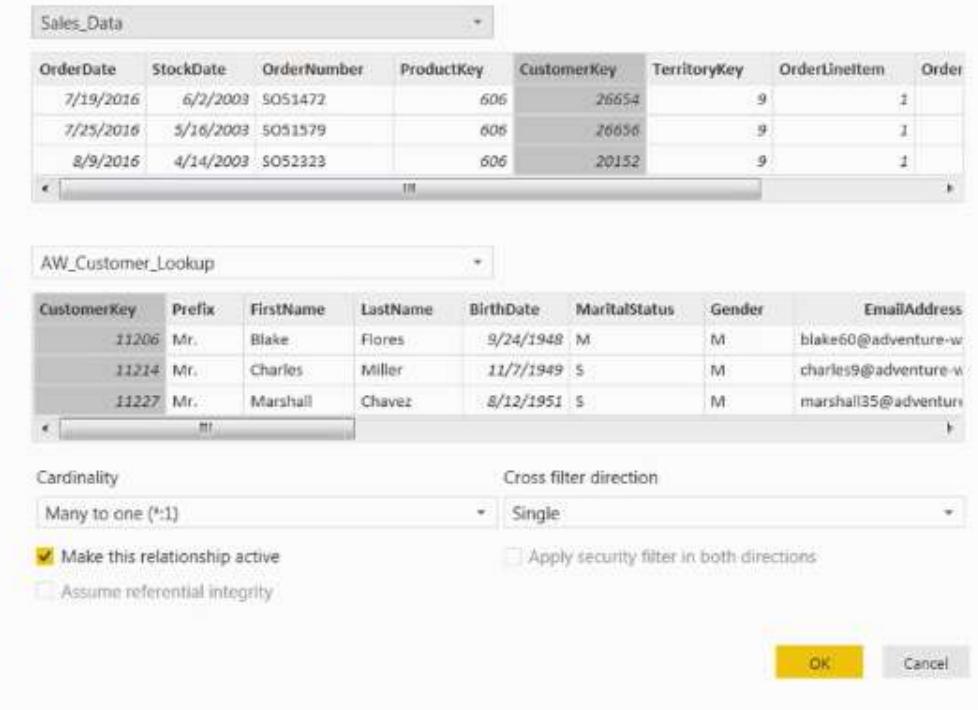
MANAGING & EDITING RELATIONSHIPS



The “**Manage Relationships**” dialog box allows you to **add, edit, or delete** table relationships

Edit relationship

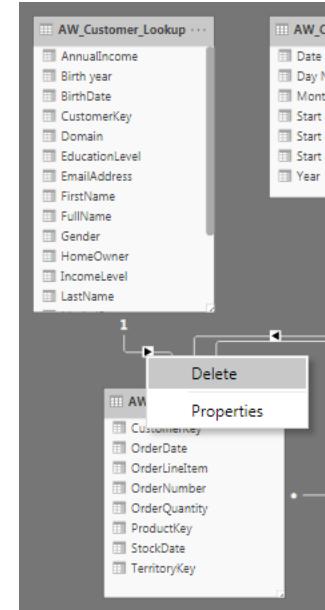
Select tables and columns that are related.



Editing tools allow you to **activate/deactivate** relationships, view **cardinality**, and modify the **cross filter direction** (stay tuned!)

Demo

- Try to use (Manage Relationship) to edit sales date relationship to and change (Order Date) to (Stock Date).
- Also we can delete the unwanted Relationship by simply clicking on it and Choose delete.



Edit relationship

Select tables and columns that are related.

AW-Sales							
OrderDate	StockDate	OrderNumber	ProductKey	CustomerKey	TerritoryKey	OrderLine	CustomerName
Sunday, July 05, 2015	Monday, June 03, 2002	SO46718	360	12570	9		
Tuesday, July 07, 2015	Monday, April 22, 2002	SO46736	360	12341	9		
Sunday, July 12, 2015	Sunday, May 05, 2002	SO46776	360	12356	9		

AW_Calendar_Lookup						
Date	Day Name	Start of Week	Start of Month	Month Name	Start of Year	Year
Friday, January 01, 2016	Friday	Monday, December 28, 2015	Friday, January 01, 2016	January	Friday, Jan	2016
Saturday, January 02, 2016	Saturday	Monday, December 28, 2015	Friday, January 01, 2016	January	Friday, Jan	2016
Sunday, January 03, 2016	Sunday	Monday, December 28, 2015	Friday, January 01, 2016	January	Friday, Jan	2016

Cardinality

Many to one (*:1)

Cross filter direction

Single

Make this relationship active

Apply security filter in both directions

Assume referential integrity

OK

Cancel

ACTIVE VS. INACTIVE RELATIONSHIPS

The screenshot shows the Power BI Data Model view. On the left, the **Calendar** table is listed with columns: Date, ∑ Day, Day Name, Start of Week, ∑ Month, Month Name, Start of Month, ∑ Year, and ∑ Day of Week. On the right, the **Sales_Data** table is listed with columns: OrderDate, StockDate, OrderNumber, ProductKey, CustomerKey, TerritoryKey, OrderLineItem, and Order. A relationship is being established between the Date column in the Calendar table and either the OrderDate or StockDate column in the Sales_Data table. The relationship line is highlighted with a yellow box. A tooltip indicates "Assume referential integrity". In the "Edit relationship" dialog, two options are shown:

- Relationship 1 (Calendar to OrderDate):** Shows the Sales_Data table with OrderDate selected. The "Make this relationship active" checkbox is checked (highlighted with a yellow box). The tooltip "Assume referential integrity" is visible.
- Relationship 2 (Calendar to StockDate):** Shows the Sales_Data table with StockDate selected. The "Make this relationship active" checkbox is unchecked (highlighted with a yellow box). The tooltip "Assume referential integrity" is visible.

Relationship Details:

- Cardinality:** Many to one (*:1)
- Cross filter direction:** Single
- Apply security filter:** Unchecked

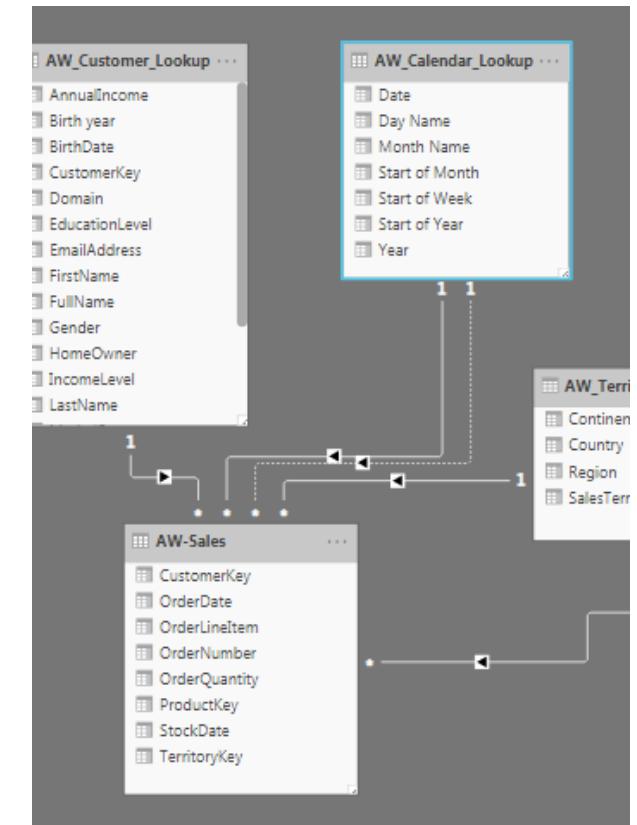
Data Preview:

OrderDate	StockDate	OrderNumber	ProductKey	CustomerKey	TerritoryKey	OrderLineItem	Order
7/23/2016	6/2/2003	S051472	606	28654			
7/25/2016	5/16/2003	S051579	606	28656			
8/9/2016	4/14/2003	S052323	606	20152			

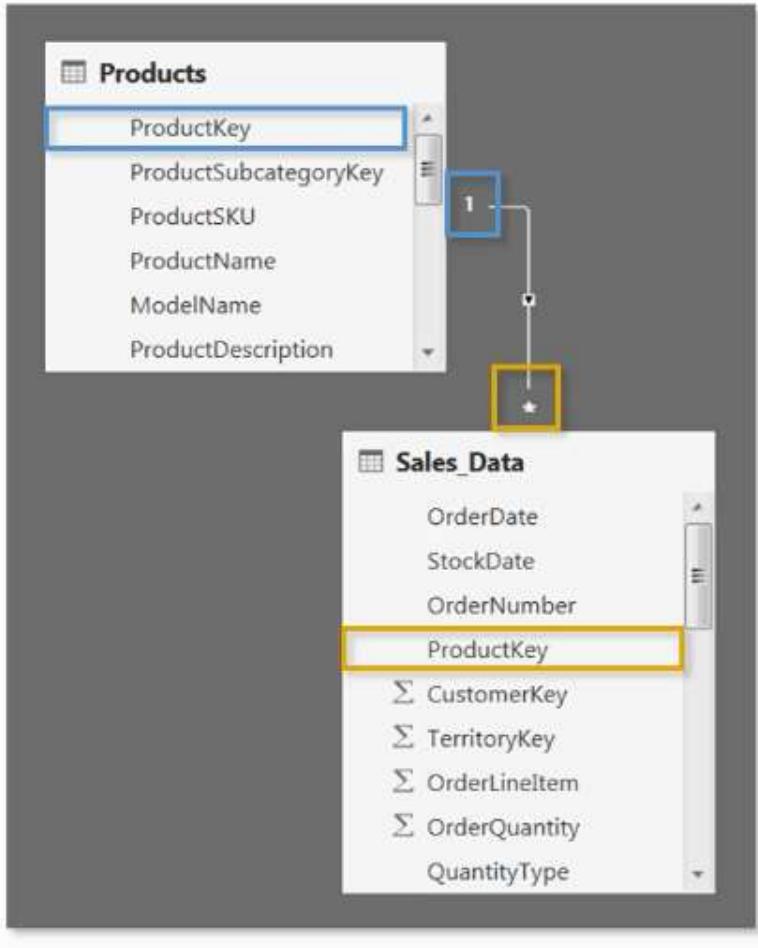
Date	Day	Day Name	Start of Week	Month	Month Name	Start
1/1/2016	1	Friday	12/27/2015	1	January	
1/2/2016	2	Saturday	12/27/2015	1	January	
1/3/2016	3	Sunday	1/1/2016	1	January	

Demo

- Add relationship between stock date (sales table) with date from (calender_lookup)
- Notice that one line is solid and other not, this is Because only one relationship can be active at a Time so by default one relation ship is active and other One inactive.



RELATIONSHIP CARDINALITY



Cardinality refers to the *uniqueness of values* in a column

- For our purposes, all relationships in the data model should follow a “**one-to-many**” cardinality; **one** instance of each *primary key*, but potentially **many** instances of each *foreign key*

*In this case, there is only **ONE instance of each ProductKey** in the Products table (noted by the “1”), since each row contains **attributes of a single product** (Name, SKU, Description, Retail Price, etc)*

*There are **MANY instances of each ProductKey** in the Sales_Data table (noted by the asterisk *), since there are **multiple sales associated with each product***

CARDINALITY CASE STUDY: MANY-TO-MANY

product_id	product_name	product_sku
4	Washington Cream Soda	64412155747
4	Washington Diet Cream Soda	81727382373
5	Washington Diet Soda	85561191439
7	Washington Diet Cola	20191444754
8	Washington Orange Juice	89770532250

date	product_id	transactions
1/1/2017	4	12
1/2/2017	4	9
1/3/2017	4	11
1/1/2017	5	16
1/2/2017	5	19
1/1/2017	7	11

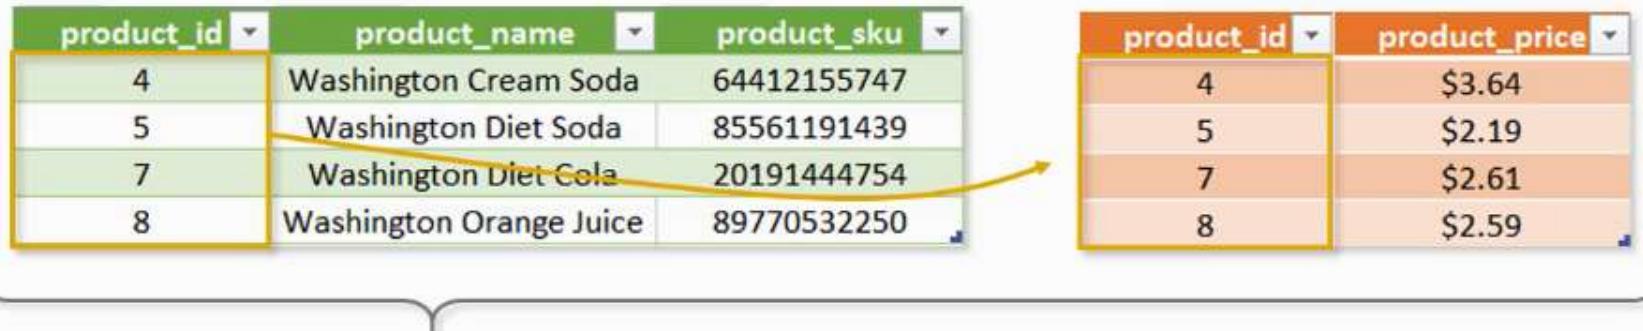
Create relationship

You can't create a relationship between these two columns because one of the columns must have unique values.

OK

- If we try to connect these tables using **product_id**, we'll get a "**many-to-many relationship**" error since there are multiple instances of each ID in both tables
- Even if we *could* create this relationship, how would you know which product was actually sold on each date – *Cream Soda* or *Diet Cream Soda*?

CARDINALITY CASE STUDY: ONE-TO-ONE



- Connecting the two tables above using the **product_id** field creates a **one-to-one relationship**, since each ID only appears once in each table
- Unlike many-to-many, there is nothing *illegal* about this relationship; it's just **inefficient**

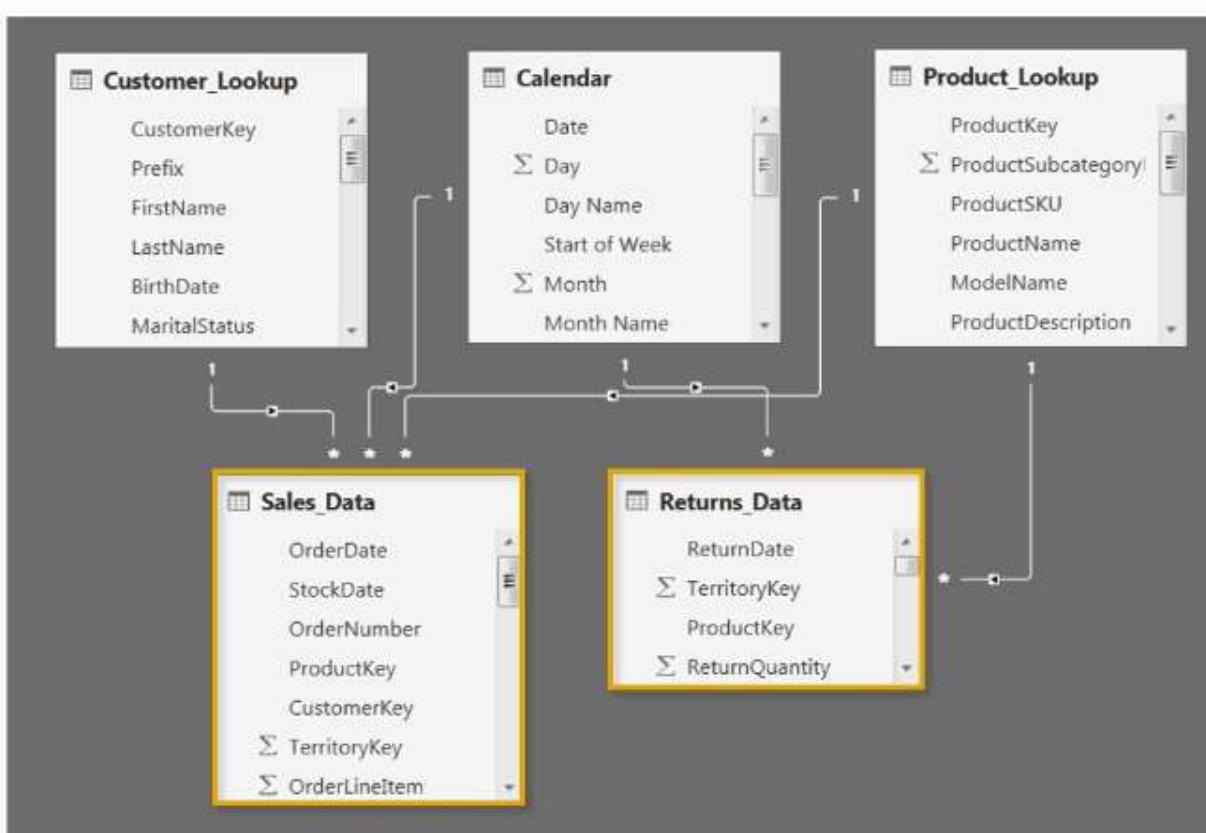
To eliminate the inefficiency, you could simply **merge the two tables** into a single, valid lookup

NOTE: this still respects the laws of normalization, since all rows are unique and capture attributes related to the primary key

The diagram shows a merged table where the two original tables have been joined together. The resulting table has four columns: product_id, product_name, product_sku, and product_price. It contains four rows with IDs 4, 5, 7, and 8, corresponding to the rows in the original tables. The row with ID 7 is highlighted with a yellow border.

product_id	product_name	product_sku	product_price
4	Washington Cream Soda	64412155747	\$3.64
5	Washington Diet Soda	85561191439	\$2.19
7	Washington Diet Cola	20191444754	\$2.61
8	Washington Orange Juice	89770532250	\$2.59

CONNECTING MULTIPLE DATA TABLES



This model contains two data tables:
Sales_Data and **Returns_Data**

- Note that the **Returns** table connects to **Calendar** and **Product_Lookup** just like the **Sales** table, but without a *CustomerKey* field it cannot be joined to **Customer_Lookup**
- This allows us to analyze sales and returns within the same view, **but only if we filter or segment the data using shared lookups**
 - In other words, we know which **product** was returned and on which **date**, but nothing about which **customer** made the return



HEY THIS IS IMPORTANT!

In general, never create direct relationships between data tables; instead, connect them through shared lookups

Demo

- Import (AW_return) table
- Form a relationship between (AW_return) and (Territories_lookup,Calende_lookup ,Product_lookup)
- Go to report view tab
- Select the previous matrix
- Drag (return quantity) to values field with (order quantity)

Next: Delete (Product column) from value field and replace it with Gender from (customer_lookup), you will notice that return quantity values is fixed for all rows as there is no relationship between (customer_lookup) and (return data)

Note: Don't form relationship between two data tables as it will be many-to-many relationship

FILTER FLOW



Here we have two data tables (**Sales_Data** and **Returns_Data**), connected to **Territory_Lookup**

Note the filter directions (shown as arrows) in each relationship; by default, **these will point from the “one” side of the relationship (lookups) to the “many” side (data)**

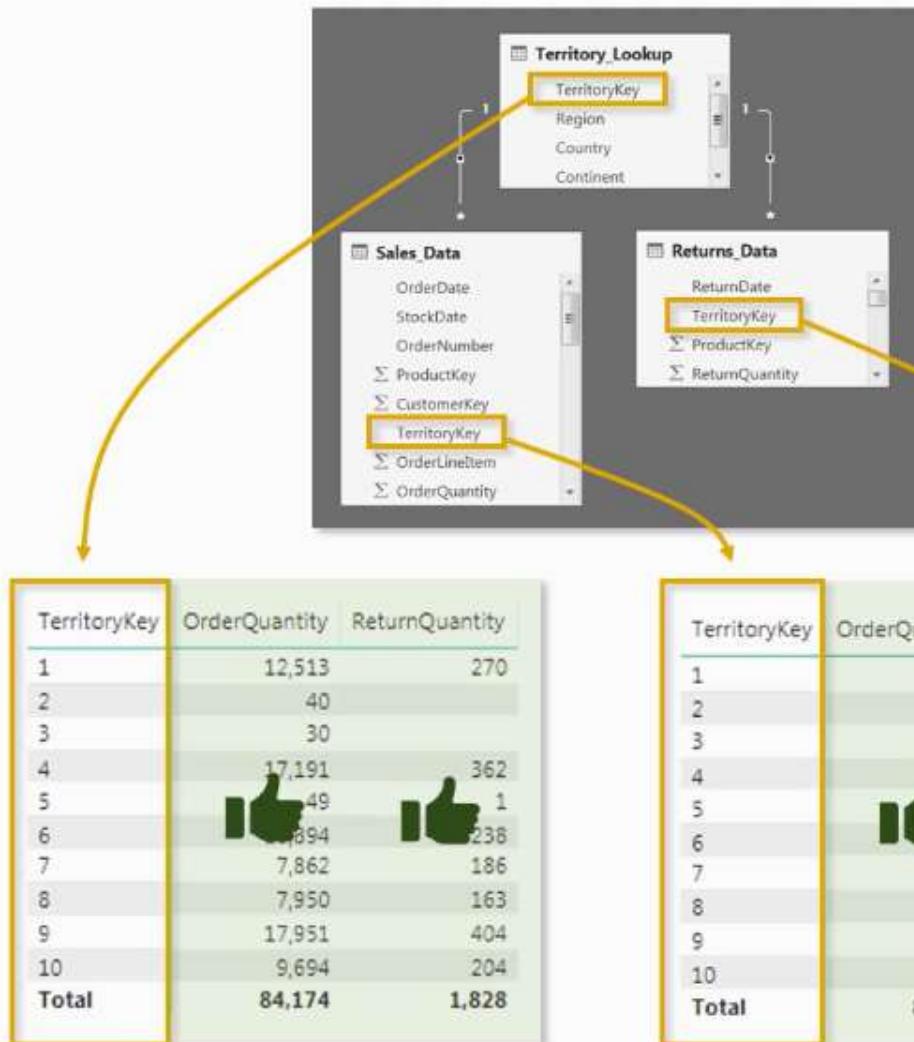
- When you filter a table, that filter context is passed along to all related “*downstream*” tables (following the direction of the arrow)
- Filters **cannot** flow “*upstream*” (against the direction of the arrow)



PRO TIP:

Arrange your lookup tables **above** your data tables in your model as a visual reminder that filters flow “*downstream*”

FILTER FLOW (CONT.)



In this case, the only valid way filter both **Sales** and **Returns** data by Territory is to use the **TerritoryKey** field from the **Territory_Lookup** table, which is upstream and related to *both* data tables

- Filtering using **TerritoryKey** from the **Sales** table yields incorrect **Returns** values, since the filter context *cannot flow upstream* to either one of the other tables
- Similarly, filtering using **TerritoryKey** from the **Returns** table yields incorrect **Sales** data; in addition, **only territories that registered returns are visible in the table** (even though they registered sales)

1) Filtering using **TerritoryKey** from the **Territory_Lookup** table

2) Filtering using **TerritoryKey** from the **Sales_Data** table

3) Filtering using **TerritoryKey** from the **Returns_Data** table

Demo

- Replace gender column from matrix to territory key from territory_lookup table, you will notice that both order quantities, and return quantities are responding as they are connected to territory_lookup table.

TWO-WAY FILTERS

Edit relationship

Select tables and columns that are related.

Sales_Data

OrderDate	StockDate	OrderNumber	ProductKey	CustomerKey	TerritoryKey	OrderLineItem	OrderQuantity
7/19/2016	6/2/2003	S051472	606	26654	9	1	
7/25/2016	5/16/2003	S051579	606	26656	9	1	
8/9/2016	4/14/2003	S052323	606	20152	9	1	

Territory_Lookup

TerritoryKey	Region	Country	Continent
1	Northwest	United States	North America
2	Northeast	United States	North America
3	Central	United States	North America

Cardinality

Many to one (*:1)

Cross filter direction

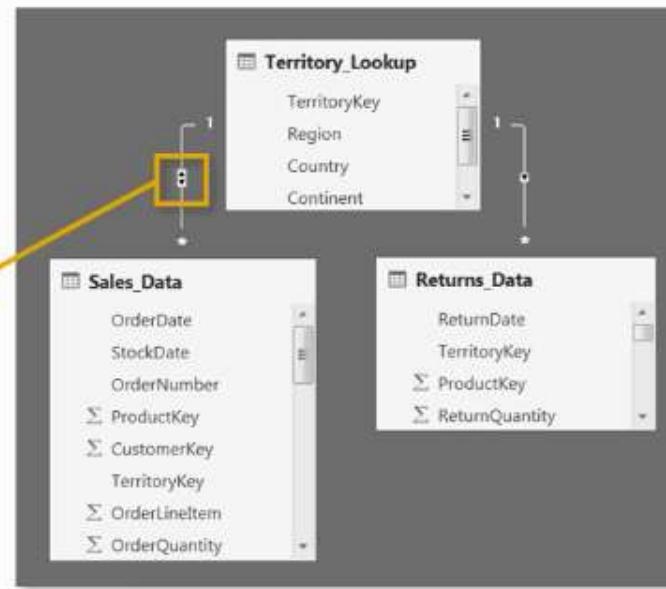
Both

Make this relationship active

Assume referential integrity

Apply security filter in both directions

OK Cancel

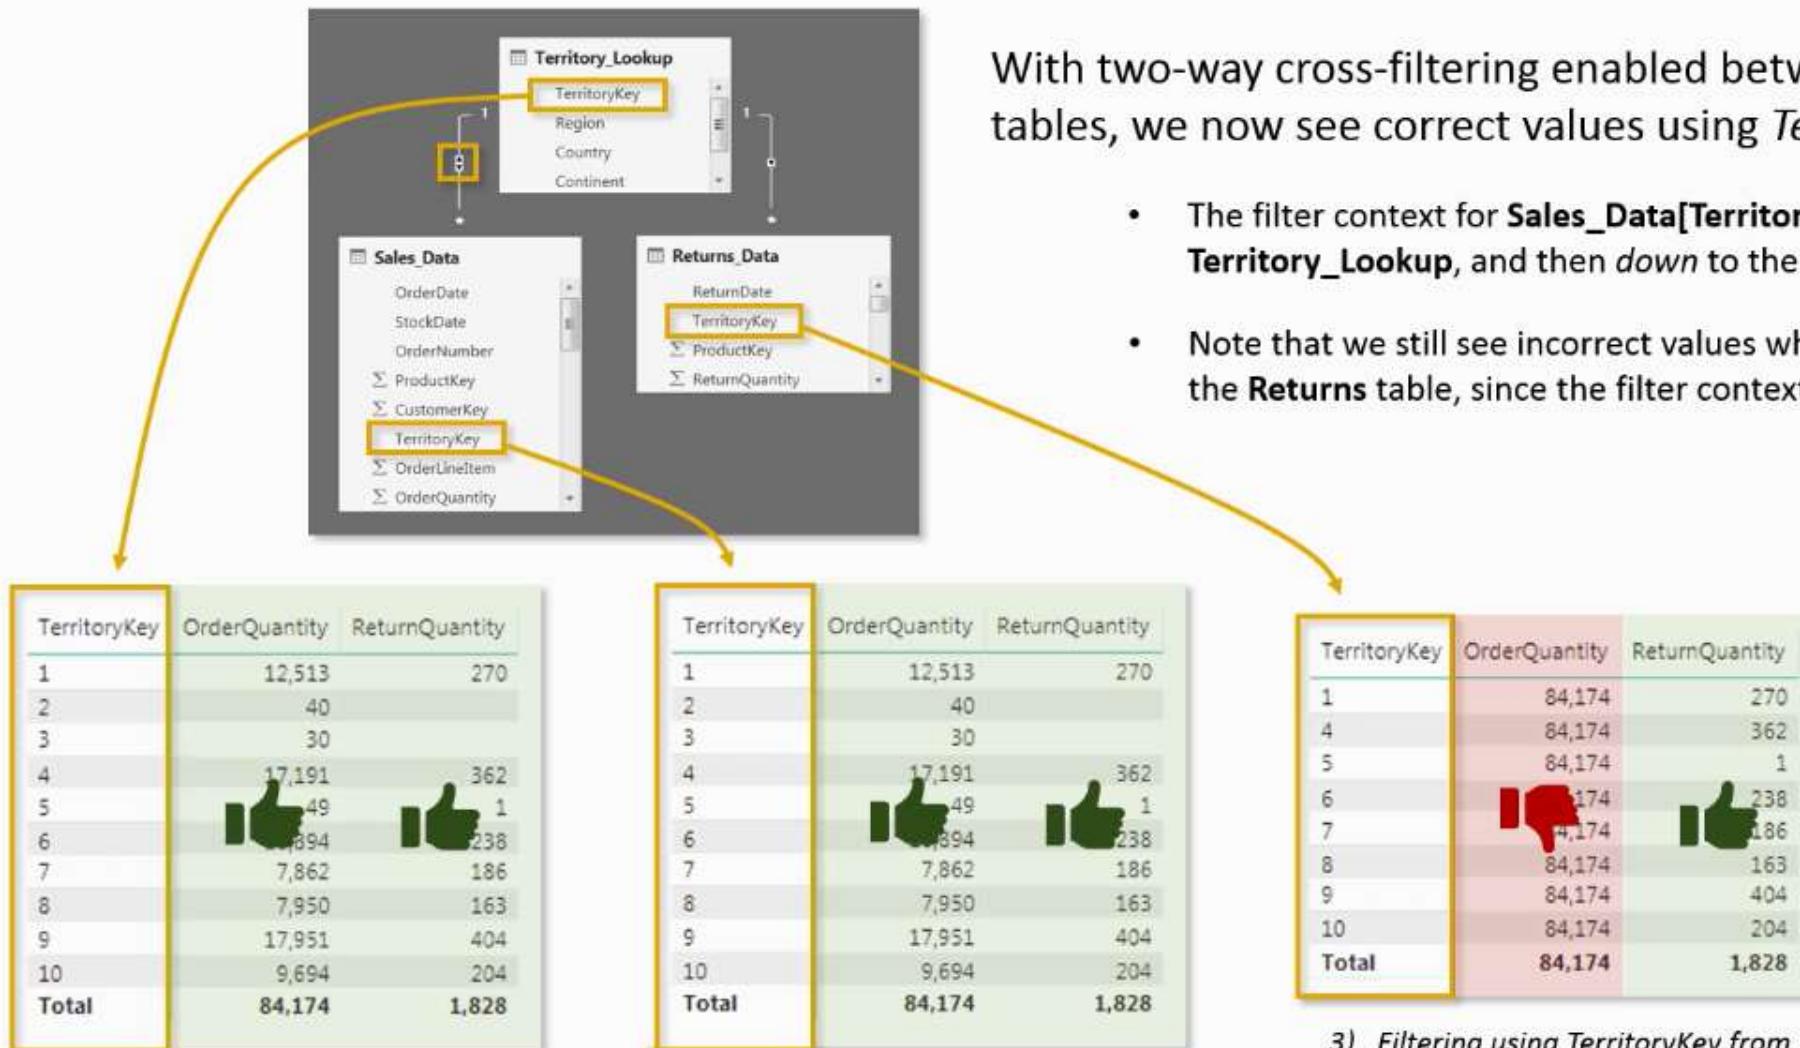


Updating the filter direction between **Sales** and **Territory** from “Single” to “Both” allows filter context to flow both ways

- This means that filters applied to the **Sales_Data** table will pass to the lookup, and then down to the **Returns_Data** table

NOTE: The “Apply security filter in both directions” option relates to row-level security (RLS) settings, which are not covered in this course

TWO-WAY FILTERS (CONT.)



With two-way cross-filtering enabled between the **Sales** and **Territory** tables, we now see correct values using *TerritoryKey* from either table

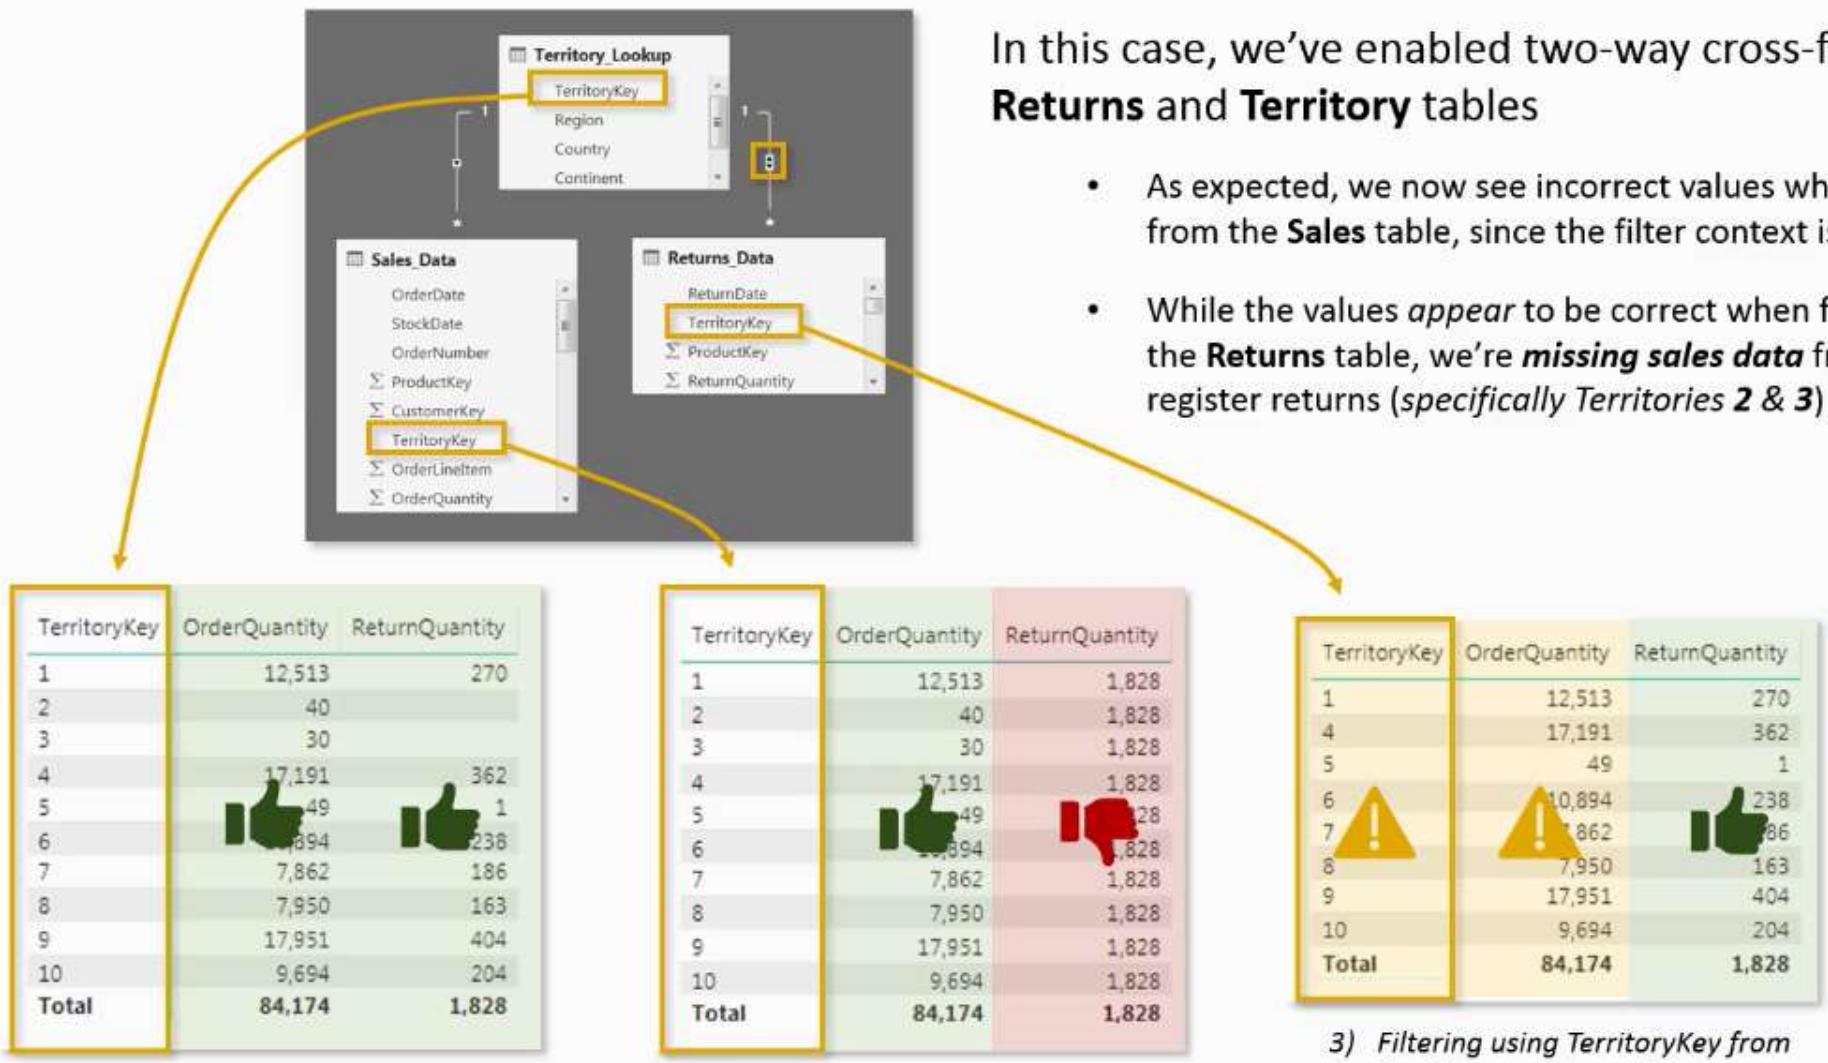
- The filter context for **Sales_Data[TerritoryKey]** now passes *up* to the **Territory_Lookup**, and then *down* to the **Returns_Data** table
- Note that we still see incorrect values when filtering using *TerritoryKey* from the **Returns** table, since the filter context is isolated to that single table

1) Filtering using TerritoryKey from the Territory_Lookup table

2) Filtering using TerritoryKey from the Sales_Data table

3) Filtering using TerritoryKey from the Returns_Data table

TWO-WAY FILTERS (CONT.)



In this case, we've enabled two-way cross-filtering between the **Returns** and **Territory** tables

- As expected, we now see incorrect values when filtering using *TerritoryKey* from the **Sales** table, since the filter context is isolated to that single table
- While the values *appear* to be correct when filtering using *TerritoryKey* from the **Returns** table, we're **missing sales data** from any territories that didn't register returns (*specifically Territories 2 & 3*)

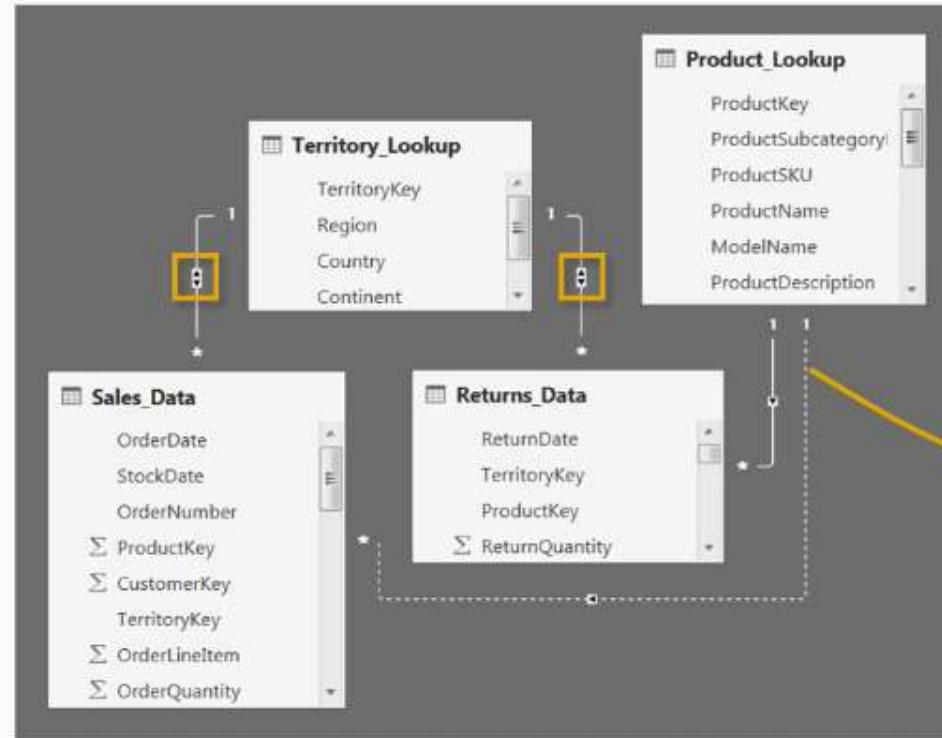
Since no information about Territory 2 or 3 is passed from the **Returns_Data** table to **Territory_Lookup**, they get filtered out of the lookup, and subsequently filtered out of the **Sales_Data**

1) Filtering using TerritoryKey from the **Territory_Lookup** table

2) Filtering using TerritoryKey from the **Sales_Data** table

3) Filtering using TerritoryKey from the **Returns_Data** table

TWO-WAY FILTERS: A WORD OF WARNING



Use two-way filters carefully, and **only when necessary***

- If you try to use multiple two-way filters in a more complex model, you run the risk of creating "**ambiguous relationships**" by introducing multiple filter paths between tables:

! You can't create a direct active relationship between **Sales_Data** and **Product_Lookup** because that would introduce ambiguity between tables **Product_Lookup** and **Territory_Lookup**. To make this relationship active, deactivate or delete one of the relationships between **Product_Lookup** and **Territory_Lookup** first.

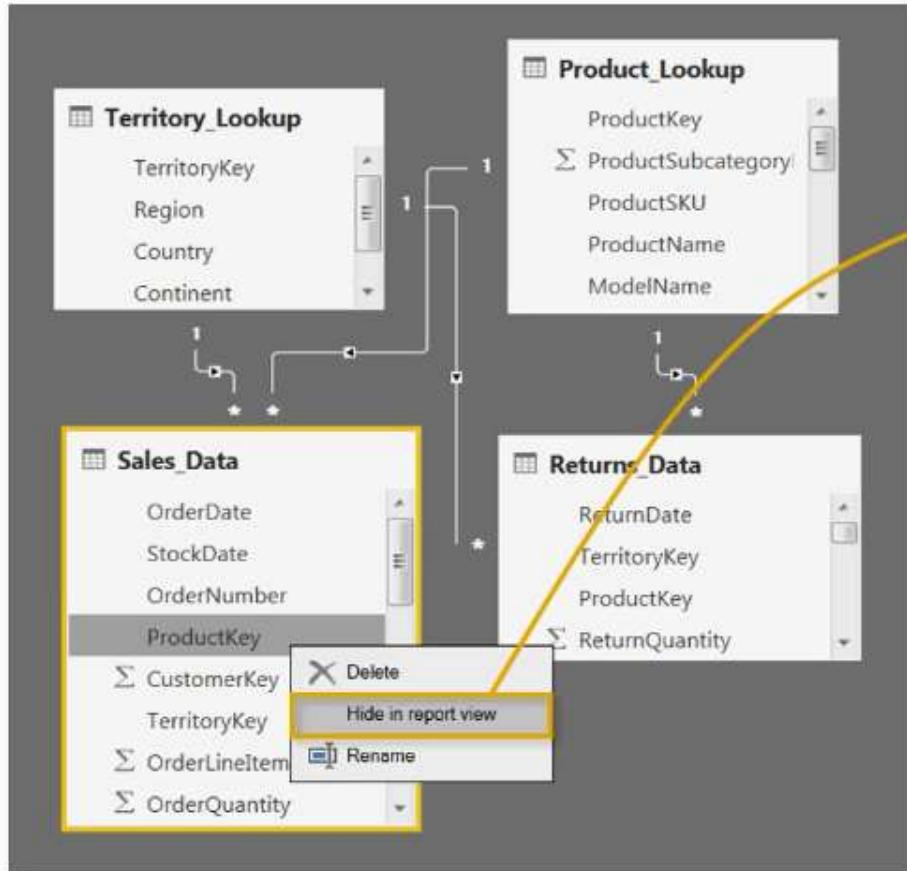
In this model, filter context from the **Product_Lookup** table can pass down to **Returns_Data** and up to **Territory_Lookup**, which would filter accordingly based on the **TerritoryKey** values passed from the **Returns** table

If we were able to activate the relationship between **Product_Lookup** and **Sales_Data** as well, filters could pass from the **Product_Lookup** table through EITHER the **Sales** or **Returns** table to reach the **Territory_Lookup**, which could yield conflicting filter context

PRO TIP:

Design your models with **one-way filters** and **1-to-Many cardinality**, unless more complex relationships are necessary

HIDING FIELDS FROM REPORT VIEW



Hiding fields from Report View makes them inaccessible from the Report tab (*although they can still be accessed within the **Data** or **Relationships** views*)

This is commonly used to prevent users from filtering using invalid fields, or to hide irrelevant metrics from view



PRO TIP:

Hide the **foreign key columns** in your data tables to force users to filter using the **primary keys** in the lookup tables

BEST PRACTICES: DATA MODELING



Focus on building a normalized model from the start

- *Make sure that each table in your model serves a single, distinct purpose*
- *Use relationships vs. merged tables; long & narrow tables are better than short & wide*



Organize lookup tables *above* data tables in the diagram view

- *This serves as a visual reminder that filters flow “downstream”*



Avoid complex cross-filtering unless absolutely necessary

- *Don’t use two-way filters when 1-way filters will get the job done*



Hide fields from report view to prevent invalid filter context

- *Recommend hiding foreign keys from data tables, so that users can only access valid fields*

Quiz

Question 1:

Why might you normalize a data model?

- To make it more efficient
- To minimize the risk of error during data modifications
- To simplify queries and analysis
- All of the above

Question 2:

Which type of table generally contains foreign keys?

Lookup Tables

Data Tables

Dimension Tables

All Tables

Question 3:

TRUE or FALSE: You can create multiple active relationships from one table to the same primary key

- True
- False

Question 4:

Which type of cardinality describes the relationships in our Adventure Works data model?

One-to-One

One-to-Many

Many-to-Many

Question 5:

Generally speaking, what's the best way to "connect" two data tables within the same model?

- Deduplicate the foreign keys in one table to create a one-to-many relationship
- Deduplicate the foreign keys in both tables
- Match the row-level granularity in order to define a one-to-one relationship
- Relate them through shared lookup tables

Question 6:

Why is it sometimes helpful to arrange lookup tables *above* data tables in the Relationships view?

- To remind you that filters flow "downstream"
- Because they generally contain the most fields
- Because it makes it easier to create new relationships
- Because lookup tables contain the fewest number of fields

Question 8:

What can you do to prevent users from filtering with invalid fields from a data model?

Politely (but firmly) ask them to cut it out

Threaten them with public humiliation

Tell them that "there will be consequences" if they do it again, then slowly back out of the room without breaking eye contact

Activate the "Hide in Report View" option

Question 9:

Which type of database schema could be described as containing a central data table surrounded by individual lookup tables?

Star

Snowflake

Snowball

Spiral

Exercise: Creating Table Relationships & Data Models in Power BI

- Using your Adventure Works report file, complete the following:
- 1) Navigate to the **RELATIONSHIPS** view, and perform the following actions
 - Right-click to delete each relationship between **AW_Sales**, **AW_Customer_Lookup** and **AW_Calendar_Lookup** (*including both date fields*)
 - Use the **Manage Relationships** tool to delete *all* remaining relationships between all tables

-
- 2) Recreate all table relationships (*using any method you prefer*), and confirm the following :
 - Cardinality is **1-to-Many** for all relationships
 - Filters are all **One-Way** (*no two-way filters*)
 - Filter direction correctly flows "**downstream**" to data tables
 - Data tables are **not connected** directly to one another
 - Both data tables are connected to **all valid lookup tables**
 - Product-related tables follow a **snowflake schema**

-
- 3) Return to the **REPORT** view, and complete the following:
 - Edit (*or insert*) the matrix visual to show **ReturnQuantity (values)** by **CategoryName (rows)** from the **AW_Product_Category_Lookup** table
 - *Which category saw the highest volume of returns? How many?*
 - Replace **CategoryName** with **Year** from the **AW_Calendar_Lookup** table
 - *How many returns do you see in 2015 vs. 2016?*
 - Replace **Year** with **FullName** from the **AW_Customer_Lookup** table
 - *What do you see, and why?*
 - Update the matrix to show both **OrderQuantity** and **ReturnQuantity (values)** by **ProductKey (rows)** from the **AW_Product_Lookup** table
 - *What was the total OrderQuantity for Product #338?*

-
- 4) Unhide the **ProductKey** field from the **AW_Retruns** tables (*using either the DATA or RELATIONSHIPS view*):
 - In the matrix,
replace **ProductKey** from **AW_Product_Lookup** with **ProductKey** from
the **AW_Retruns** table
 - *Why do we see the same repeating values for OrderQuantity?*
 - Edit the relationship between **AW_Retruns** and **AW_Product_Lookup** to
change the cross filter direction from *Single* to *Both*
 - *Why does the visual now show OrderQuantity values by product, even though we are using ProductKey from AW_Retruns?*
 - *How many orders do we see now for Product #338? What's going on here?*

-
- 5) Complete the following:
 - Change the cross filter direction between **AW_RetURNS** and **AW_Product_Lookup** back to *single (One-Way)*
 - Hide the **ProductKey** field in the **AW_RetURNS** table from report view (*and any other foreign keys, if necessary*)
 - Update the matrix to show **ProductKey** from the **AW_Product_Lookup**, rather than **AW_RetURNS**
 - **Recommendation:** Save a separate backup copy of the .pbix file (*i.e. "AdventureWorks_Report_Backup"*)



Section 4: Analyzing Data with DAX Calculations in Power BI

MEET DAX

Data Analysis Expressions, commonly known as **DAX**, is the formula language that drives Power BI. With DAX, you can:

- Add **calculated columns** and **measures** to your model, using intuitive syntax
- Go beyond the capabilities of traditional “grid-style” formulas, with powerful and flexible functions built specifically to work with relational data models

Two ways to use DAX

The image displays two screenshots illustrating the use of DAX. The left screenshot shows the Power BI Data View with a table of data and a calculated column named 'Parent' highlighted. The right screenshot shows the Power BI Model view with a context menu open over a table, demonstrating how to create new measures like 'Total Orders' and 'Total Revenue'.

1) Calculated Columns

EducationLevel	Occupation	HomeOwner	Full Name	User Name	Domain	IncomeLevel	Parent
Partial College	Professional	Y	Mrs. Bala Flores	bala60	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Charles Miller	charles9	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Marshall Chavet	marshall35	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Levi Chandra	levi1	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Sean Allen	sean49	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. James Walker	james96	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Cameron Yang	cameron23	Adventure Works	Average	Yes
Partial College	Professional	N	Mrs. Keith Raje	keith17	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Richard Coleman	richard61	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Robert Lewis	robert81	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Jonathan Polson	jonathan73	Adventure Works	Average	Yes
Partial College	Professional	Y	Mrs. Robert Wang	robert99	Adventure Works	Average	Yes
Partial College	Professional	N	Mrs. Angel Scott	angel38	Adventure Works	Average	Yes

2) Measures

```
Total Orders = DISTINCTCOUNT(Sales_Data[OrderNumber])
Total Revenue = SUMX(Sales_Data, Sales_Data[OrderQuantity] * RELATED(Product_Lookup[ProductPrice]))
Quantity Ordered = SUM(Sales_Data[OrderQuantity])
```

CALCULATED COLUMNS

Calculated columns allow you to add new, formula-based columns to tables

- No “A1-style” references; calculated columns refer to **entire tables or columns**
- Calculated columns generate values for each row, which are **visible within tables in the Data view**
- Calculated columns understand **row context**; they’re great for defining properties based on information in each row, but generally useless for aggregation (*SUM, COUNT, etc*)



HEY THIS IS IMPORTANT!

As a rule of thumb, use calculated columns when you want to “stamp” static, fixed values to each row in a table (*or use the Query Editor!*)

DO NOT use calculated columns for aggregation formulas, or to calculate fields for the “Values” area of a visualization (use **measures** instead)



PRO TIP:

Calculated columns are typically used for filtering data, rather than creating numerical values

CALCULATED COLUMNS (EXAMPLES)

The screenshot shows a Power BI Data View window. A calculated column named "Parent" is defined with the formula: `Parent = IF(Customer_Lookup[TotalChildren]>0, "Yes", "No")`. The "Parent" column is highlighted with a yellow border and contains the value "Yes" for all rows. A green thumbs-up icon is overlaid on the column header. To the right, a "FIELDS" pane lists various tables and columns, with "Customer_Lookup" expanded to show its fields.

level	Occupation	HomeOwner	Full Name	User Name	Domain	IncomeLevel	Parent
College	Professional	Y	Mr. Blake Flores	blake60	Adventure Works	Average	Yes
College	Professional	Y	Mr. Charles Miller	charles9	Adventure Works	Average	Yes
College	Professional	Y	Mr. Marshall Chavez	marshall35	Adventure Works	Average	Yes
College	Professional	Y	Mr. Levi Chandra	levi1	Adventure Works	Average	Yes
College	Professional	Y	Mr. Sean Allen	sean49	Adventure Works	Average	Yes
College	Professional	Y	Mr. James Walker	james96	Adventure Works	Average	Yes
College	Professional	Y	Mr. Cameron Yang	cameron23	Adventure Works	Average	Yes
College	Professional	N	Mr. Keith Raje	keith17	Adventure Works	Average	Yes
College	Professional	Y	Mr. Richard Coleman	richard61	Adventure Works	Average	Yes
College	Professional	Y	Mr. Robert Lewis	robert81	Adventure Works	Average	Yes
College	Professional	Y	Mr. Jonathan Robinson	jonathan72	Adventure Works	Average	Yes
College	Professional	Y	Mr. Robert Wang	robert36	Adventure Works	Average	Yes

In this case we've added a **calculated column** named "**Parent**", which equals "**Yes**" if the [TotalChildren] field is greater than 0, and "**No**" otherwise (*just like Excel!*)

- Since calculated columns understand **row context**, a new value is calculated in each row based on the value in the [TotalChildren] column
- This is a **valid use** of calculated columns; it creates a new row "property" that we can now use to filter or segment any related data within the model

Here we're using an aggregation function (SUM) to calculate a new column named **TotalQuantity**

- Since calculated columns do not understand **filter context**, the same grand total is returned in *every single row* of the table
- This is **not a valid use** of calculated columns; these values are statically "stamped" onto the table and can't be filtered, sliced, subdivided, etc.

The screenshot shows a Power BI Data View window. A calculated column named "TotalQuantity" is defined with the formula: `TotalQuantity = SUM(AW_Sales_Data[OrderQuantity])`. The "TotalQuantity" column is highlighted with a yellow border and contains the value 84174 for every row. A red thumbs-down icon is overlaid on the column header. To the right, a "FIELDS" pane lists various tables and columns, with "AW_Sales_Data" expanded to show its fields.

OrderDate	OrderNumber	ProductKey	CustomerKey	TerritoryKey	OrderLineItem	OrderQuantity	QuantityType	TotalQuantity
2003-01-01	SO46718	360	12570	5	1	1	Single Item	84174
2002-02-01	SO46736	360	12343	5	1	1	Single Item	84174
2002-03-01	SO46776	360	12356	5	1	1	Single Item	84174
2002-04-01	SO46808	360	12347	5	1	1	Single Item	84174
2002-05-01	SO46826	360	12573	5	1	1	Single Item	84174
2002-06-01	SO47075	360	12683	5	1	1	Single Item	84174
2002-07-01	SO47098	360	12657	5	1	1	Single Item	84174
2002-08-01	SO47149	360	12689	5	1	1	Single Item	84174
2002-09-01	SO47212	360	12580	5	1	1	Single Item	84174
2002-10-01	SO47302	360	12670	5	1	1	Single Item	84174
2002-11-01	SO47328	360	12681	5	1	1	Single Item	84174
2002-12-01	SO47346	360	12585	5	1	1	Single Item	84174
2002-01-01	SO47744	360	12989	5	1	1	Single Item	84174
2002-02-01	SO47745	360	12998	5	1	1	Single Item	84174

Demo

- Go to data view
- Select AW_sales table → create new column (QuantityType)
- Write the following DAX formula.

```
QuantityType = IF(AW_Sales[OrderQuantity]>1,"Multiple Items","Single Item")
```

- Select AW_sales table → create new column (TotalQuantity)
- Write the following DAX aggregated formula

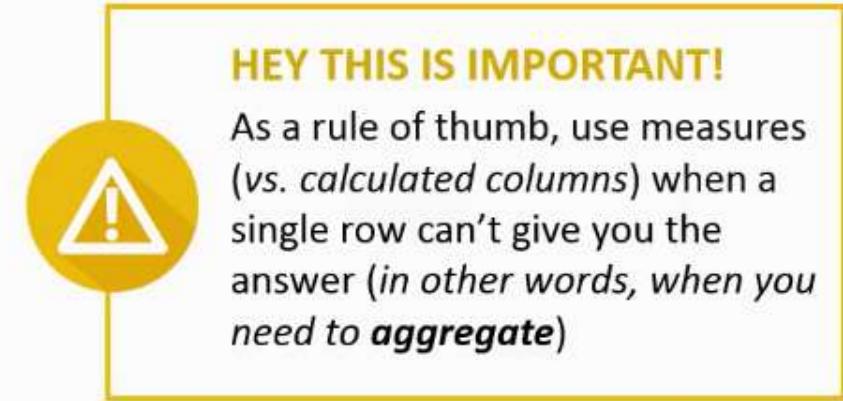
```
TotalQuantity = SUM(AW_Sales[OrderQuantity])
```

Note: Notice that we got same value repeated in each row.

MEASURES

Measures are DAX formulas used to generate new calculated values

- Like calculated columns, measures reference **entire tables** or **columns** (*no A1-style or “grid” references*)
- *Unlike calculated columns, measure values aren’t visible within tables; they can only be “seen” within a visualization like a chart or matrix (similar to a calculated field in an Excel pivot)*
- Measures are evaluated based on **filter context**, which means they recalculate when the fields or filters around them change (*like when new row or column labels are pulled into a matrix or when new filters are applied to a report*)



PRO TIP:

*Use measures to create **numerical, calculated values** that can be analyzed in the “**values**” field of a report visual*

RECAP: CALCULATED COLUMNS VS. MEASURES

CALCULATED COLUMNS

- Values are calculated based on information from each row of a table (**has row context**)
- Appends static values to each row in a table and stores them in the model (*which increases file size*)
- Recalculate on data source refresh or when changes are made to component columns
- Primarily used as **rows, columns, slicers or filters**

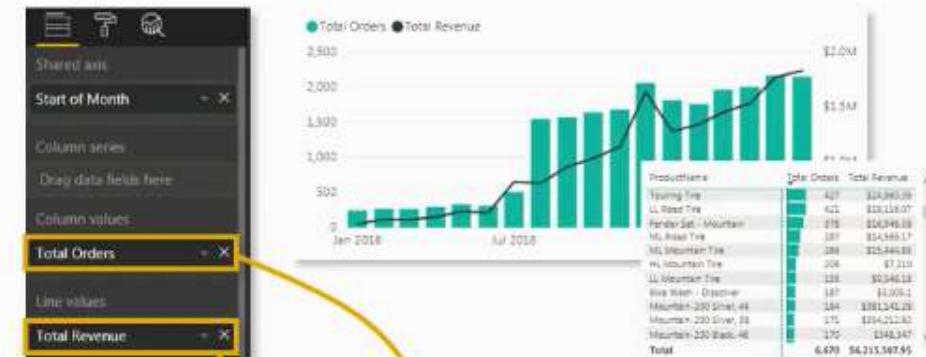
A screenshot of a Power BI data view. A table is displayed with columns: Level, Occupation, HomeOwner, Full Name, User Name, Domain, IncomeLevel, and Parent. The 'Parent' column contains values like 'Yes' and 'No'. A yellow box highlights the 'Parent' column. The status bar at the bottom shows 'Parent = IF([Customer_Lookup[TotalChildren]>0, "Yes", "No"])

Level	Occupation	HomeOwner	Full Name	User Name	Domain	IncomeLevel	Parent
Beginner	Professional	Y	Mr. Blake Flores	blake02	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. Charles Miller	charles03	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. Marshall Chavez	marshall05	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. Levi Chandra	levi11	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. Sean Allen	sean45	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. James Walker	james08	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. Cameron Yang	cameron23	Adventure Works	Average	Yes
Beginner	Professional	N	Mr. Keith Raye	keith17	Adventure Works	Average	No
Beginner	Professional	Y	Mr. Richard Coleman	richard61	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. Robert Lewis	robert01	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. Jonathan Robinson	jonthomas72	Adventure Works	Average	Yes
Beginner	Professional	Y	Mr. Robert Wang	robert36	Adventure Works	Average	Yes

Calculated columns “live” in tables

MEASURES

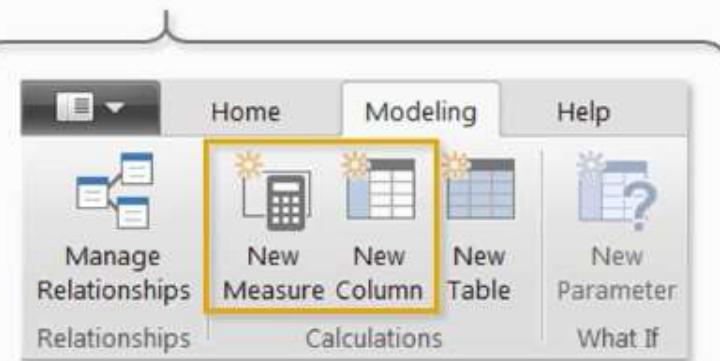
- Values are calculated based on information from any filters in the report (**has filter context**)
- Does not create new data in the tables themselves (*doesn't increase file size*)
- Recalculate in response to any change to filters within the report
- Almost *always* used within the **values** field of a visual



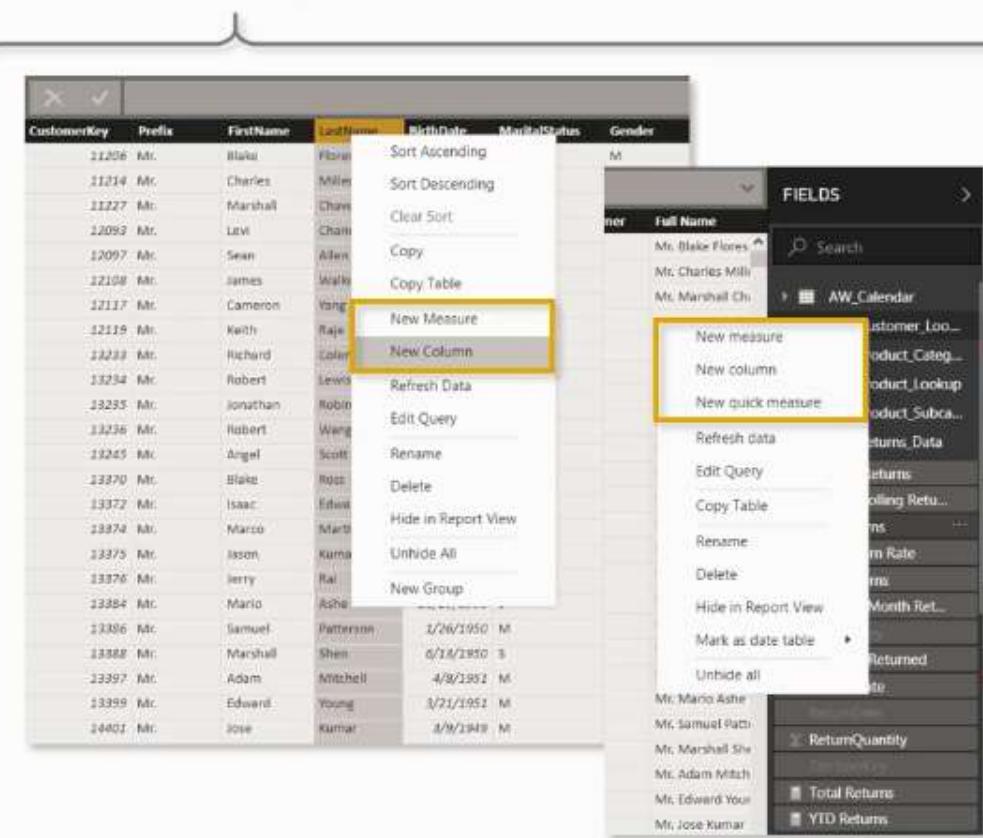
Measures “live” in visuals

ADDING COLUMNS & MEASURES

Option 1: Select “New Measure” or “New Column” from the **Modeling** tab



Option 2: Right-click within the **table** (in the **Data** view) or the **Field List** (in either the **Data** or **Report** view)



When you insert Columns or Measures using the **Modeling** tab (Option 1), they are assigned to whichever table is *currently selected*, or the *first table in the field list* by default

- Measures can be reassigned to new “Home” tables (under the “Properties” options in the **Modeling** tab), but the Option 2 allows you to be more deliberate about placing them
 - *Note: Assigning measures to specific tables doesn't have any impact on functionality – it's just a way to keep them organized*

QUICK MEASURES

The screenshot shows the 'Quick measures' dialog box in Power BI. On the left, a vertical menu lists options like 'New measure', 'New column', and 'New quick measure'. A yellow arrow points from the 'New quick measure' option to the dialog. The dialog itself has two main sections: 'Calculation' and 'Fields'. The 'Calculation' section is expanded, showing a dropdown menu with 'Average per category' selected. Below it is a list of other calculation templates: 'Aggregate per category', 'Variance per category', 'Max per category', 'Min per category', 'Weighted average per category', 'Filters', 'Time intelligence', and 'Rolling average'. The 'Fields' section shows a list of available fields from the data model, starting with 'AW_Calendar' and including 'AW_Customer_Lookup', 'AW_Product_Category_Lookup' (which is highlighted in blue), 'AW_Product_Lookup', 'AW_Product_Subcategory_Lookup', 'AW_Returns_Data', 'AW_Sales_Data', 'AW_Territory_Lookup', and 'Parameter'. At the bottom of the dialog, there's a note: 'Don't see the calculation you want? Post an idea.' followed by 'OK' and 'Cancel' buttons.

Quick Measures are pre-built formula templates that allow you to drag and drop fields, rather than write DAX from scratch

While these tools can be helpful for defining more complex measures (*like weighted averages or time intelligence formulas*), they encourage laziness and don't help you understand the fundamentals of DAX



PRO TIP:

Just say "**NO**" to quick measures
(you're better than that)

Demo

- Go to report view tab
- Select AW_salse and then write click then
Select new measure
- Write the following DAX

```
Quantity Sold = SUM(AW_Sales[OrderQuantity])
```

- Select the new quantity sold field under
AW_Sales and drag it to value.

ProductKey	OrderQuantity
214	2099
215	1940
220	1995
223	4151
226	392
229	408
232	424
235	381
310	169
311	139
312	179
313	168
314	157
320	65
322	39
324	72
Total	84174

STEP-BY-STEP MEASURE CALCULATION

CategoryName	Total Returns
Accessories	1,115
Bikes	342
Clothing	267

How exactly is this measure calculated?

- REMEMBER: This all happens *instantly* behind the scenes, every time the filter context changes

STEP 1

Filter context is detected & applied



CategoryName	Total Returns
Accessories	1,115
Bikes	342
Clothing	267

Product[CategoryName] = "Accessories"

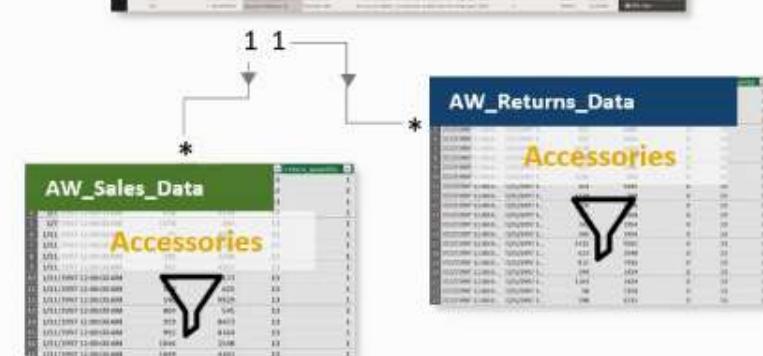
Product Table	
Accessories	Filter icon

STEP 2

Filters flow “downstream” to all related tables



Product Table	
Accessories	Filter icon



STEP 3

Measure formula evaluates against the filtered table



X ✓ Total Returns = COUNTROWS(AW_Returns_Data) ▾

Count of rows in the AW_Returns_Data table, filtered down to only rows where the product category is "Accessories"

= 1,115

DAX SYNTAX

MEASURE NAME

- Note:** Measures are always surrounded in brackets (i.e. [Total Quantity]) when referenced in formulas, so spaces are OK

Total Quantity: =**SUM(Transactions[quantity])**

FUNCTION NAME

- Calculated columns don't always use functions, but measures do:
 - In a **Calculated Column**, =Transactions[quantity] returns the value from the quantity column in each row (*since it evaluates one row at a time*)
 - In a **Measure**, =Transactions[quantity] will return an **error** since Power BI doesn't know how to translate that as a single value (*you need some sort of aggregation*)

Referenced
TABLE NAME

Referenced
COLUMN NAME

Note: This is a "fully qualified" column, since it's preceded by the table name -- table names with spaces must be surrounded by **single quotes**:

- Without a space: Transactions[quantity]
- With a space: 'Transactions Table'[quantity]

PRO TIP:

For **column** references, use the fully qualified name (i.e. **Table[Column]**)
For **measure** references, just use the measure name (i.e. **[Measure]**)



DAX OPERATORS

Arithmetic Operator	Meaning	Example
+	Addition	2 + 7
-	Subtraction	5 – 3
*	Multiplication	2 * 6
/	Division	4 / 2
^	Exponent	2 ^ 5

Comparison Operator	Meaning	Example
=	Equal to	[City] = "Boston"
>	Greater than	[Quantity] > 10
<	Less than	[Quantity] < 10
>=	Greater than or equal to	[Unit_Price] >= 2.5
<=	Less than or equal to	[Unit_Price] <= 2.5
<>	Not equal to	[Country] <> "Mexico"

Text/Logical Operator	Meaning	Example
&	Concatenates two values to produce one text string	[City] & " " & [State]
&&	Create an AND condition between two logical expressions	([State] = "MA") && ([Quantity] > 10)
(double pipe)	Create an OR condition between two logical expressions	([State] = "MA") ([State] = "CT")
IN	Creates a logical OR condition based on a given list (using curly brackets)	'Store Lookup'[State] IN { "MA", "CT", "NY" }

DAX OPERATORS

Arithmetic Operator	Meaning	Example
+	Addition	$2 + 7$
-	Subtraction	$5 - 3$
*	Multiplication	$2 * 6$
/	Division	$4 / 2$
\wedge	Exponent	$2 \wedge 5$

Pay attention to these!

Comparison Operator	Meaning	Example
=	Equal to	[City] = "Boston"
>	Greater than	[Quantity] > 10
<	Less than	[Quantity] < 10
\geq	Greater than or equal to	[Unit_Price] \geq 2.5
\leq	Less than or equal to	[Unit_Price] \leq 2.5
\neq	Not equal to	[Country] \neq "Mexico"

Text/Logical Operator	Meaning	Example
&	Concatenates two values to produce one text string	[City] & " " & [State]
&&	Create an AND condition between two logical expressions	([State] = "MA") && ([Quantity] > 10)
(double pipe)	Create an OR condition between two logical expressions	([State] = "MA") ([State] = "CT")
IN	Creates a logical OR condition based on a given list (using curly brackets)	'Store Lookup'[State] IN { "MA", "CT", "NY" }

COMMON FUNCTION CATEGORIES

MATH & STATS Functions

Basic aggregation functions as well as "iterators" evaluated at the row-level

Common Examples:

- SUM
- AVERAGE
- MAX/MIN
- DIVIDE
- COUNT/COUNTA
- COUNTROWS
- DISTINCTCOUNT

Iterator Functions:

- SUMX
- AVERAGEX
- MAXX/MINX
- RANKX
- COUNTX

LOGICAL Functions

Functions for returning information about values in a given conditional expression

Common Examples:

- IF
- IFERROR
- AND
- OR
- NOT
- SWITCH
- TRUE
- FALSE

TEXT Functions

Functions to manipulate text strings or control formats for dates, times or numbers

Common Examples:

- CONCATENATE
- FORMAT
- LEFT/MID/RIGHT
- UPPER/LOWER
- PROPER
- LEN
- SEARCH/FIND
- REPLACE
- REPT
- SUBSTITUTE
- TRIM
- UNICHAR

FILTER Functions

Lookup functions based on related tables and filtering functions for dynamic calculations

Common Examples:

- CALCULATE
- FILTER
- ALL
- ALLEXCEPT
- RELATED
- RELATEDTABLE
- DISTINCT
- VALUES
- EARLIER/EARLIEST
- HASONEVALUE
- HASONEFILTER
- ISFILTERED
- USERELATIONSHIP

DATE & TIME Functions

Basic date and time functions as well as advanced time intelligence operations

Common Examples:

- DATEDIFF
- YEARFRAC
- YEAR/MONTH/DAY
- HOUR/MINUTE/SECOND
- TODAY/NOW
- WEEKDAY/WEEKNUM

Time Intelligence Functions:

- DATESYTD
- DATESQTD
- DATESMTD
- DATEADD
- DATESINPERIOD

BASIC DATE & TIME FUNCTIONS

**DAY/MONTH/
YEAR()**

Returns the day of the month (1-31), month of the year (1-12), or year of a given date

=**DAY/MONTH/YEAR**(Date)

**HOUR/MINUTE/
SECOND()**

Returns the hour (0-23), minute (0-59), or second (0-59) of a given datetime value

=**HOUR/MINUTE/SECOND**(Datetime)

TODAY/NOW()

Returns the current date or exact time

=**TODAY/NOW**()

**WEEKDAY/
WEEKNUM()**

Returns a weekday number from 1 (Sunday) to 7 (Saturday), or the week # of the year

=**WEEKDAY/WEEKNUM**(Date, [ReturnType])

EOMONTH()

Returns the date of the last day of the month, +/- a specified number of months

=**EOMONTH**(StartDate, Months)

DATEDIFF()

Returns the difference between two dates, based on a selected interval

=**DATEDIFF**(Date1, Date2, Interval)

Demo

- Go to table view → select calendar table
- Create new column
- Write DAX formula

```
Day of week = WEEKDAY(AW_Calendar_Lookup[Date],1)
```

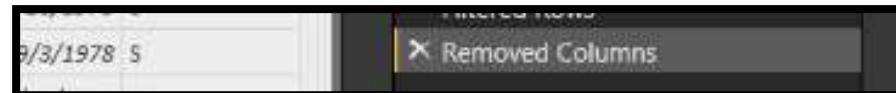
- Go to (Customer_Lookup) table (Birth Year)column , create same (birth year) column using DAX and remove the original.

```
BirthYear_CC = year(AW_Customer_Lookup[BirthDate])
```

- Also we can get same column values without using DAX functions by

```
BirthYear_CC = AW_Customer_Lookup[BirthDate].[Year]
```

- Delete (BirthYear) column generate by Query Editor and see if this step will be added.



Demo (Cont) get the age of the customers

- Create new column
- Write DAX formula to calculate the (customer age) using

```
Current Age = DATEDIFF(AW_Customer_Lookup[BirthDate], TODAY(), YEAR)
```

BASIC LOGICAL FUNCTIONS (IF/AND/OR)

IF()

Checks if a given condition is met, and returns one value if the condition is TRUE, and another if the condition is FALSE

=**IF**(LogicalTest, ResultIfTrue, [ResultIfFalse])

IFERROR()

Evaluates an expression and returns a specified value if the expression returns an error, otherwise returns the expression itself

=**IFERROR**(Value, ValueIfError)

AND()

Checks whether both arguments are TRUE, and returns TRUE if both arguments are TRUE, otherwise returns FALSE

=**AND**(Logical1, Logical2)

Note: Use the && and || operators if you want to include more than two conditions!

OR()

Checks whether one of the arguments is TRUE to return TRUE, and returns FALSE if both arguments are FALSE

=**OR**(Logical1, Logical2)

Demo, check whether the customer is parent or not?

- Go to (AW_customer_lookup) , create new column
- Write DAX formula to check whether a customer is parent or not

```
parent = if(AW_Customer_Lookup[TotalChildren]>0,"Yes","No")
```

Demo , Identify weekend day from Day of Week column

- Go to (AW_Calender_lookup)
- Write the following to check whether the day is week end or not

```
weekend = if(or(AW_Calendar_Lookup[Day of week]=6,AW_Calendar_Lookup[Day of week]=7),"weekend","No")
```

TEXT FUNCTIONS

LEN()

Returns the number of characters in a string

=**LEN**(Text)

Note: Use the & operator as a shortcut, or to combine more than two strings!

CONCATENATE()

Joins two text strings into one

=**CONCATENATE**(Text1, Text2)

**LEFT/MID/
RIGHT()**

Returns a number of characters from the start/middle/end of a text string

=**LEFT/RIGHT**(Text, [NumChars])

=**MID**(Text, StartPosition, NumChars)

**UPPER/LOWER/
PROPER()**

Converts letters in a string to upper/lower/proper case

=**UPPER/LOWER/PROPER**(Text)

SUBSTITUTE()

Replaces an instance of existing text with new text in a string

=**SUBSTITUTE**(Text, OldText, NewText, [InstanceNumber])

SEARCH()

Returns the position where a specified string or character is found, reading left to right

=**SEARCH**(FindText, WithinText, [StartPosition], [NotFoundValue])

Demo

- Go to AW_Csutomer_lookup
- Create new column
- Write DAX formula to Concatenate (prefix , First name, last name)

```
fullname_CC = AW_Customer_Lookup[Prefix] &
" " & AW_Customer_Lookup[FirstName] & " " &
AW_Customer_Lookup[LastName]
```

Demo

- Go to (AW_Calendar_lookup) table
- Write DAX formula to get a shortcut for the month name (first 3 char of each month)

```
month name _CC = left(AW_Calendar_Lookup[Month Name],3)
```

- Change the values in (month name) column to upper case

```
month name _CC = upper(left(AW_Calendar_Lookup[Month Name],3))
```

BASIC MATH & STATS FUNCTIONS

SUM()

Evaluates the sum of a column

=**SUM**(ColumnName)

AVERAGE()

Returns the average (arithmetic mean) of all the numbers in a column

=**AVERAGE**(ColumnName)

MAX()

Returns the largest value in a column or between two scalar expressions

=**MAX**(ColumnName) or =**MAX**(Scalar1, [Scalar2])

MIN()

Returns the smallest value in a column or between two scalar expressions

=**MIN**(ColumnName) or =**MIN**(Scalar1, [Scalar2])

DIVIDE()

Performs division and returns the alternate result (or blank) if div/0

=**DIVIDE**(Numerator, Denominator, [AlternateResult])

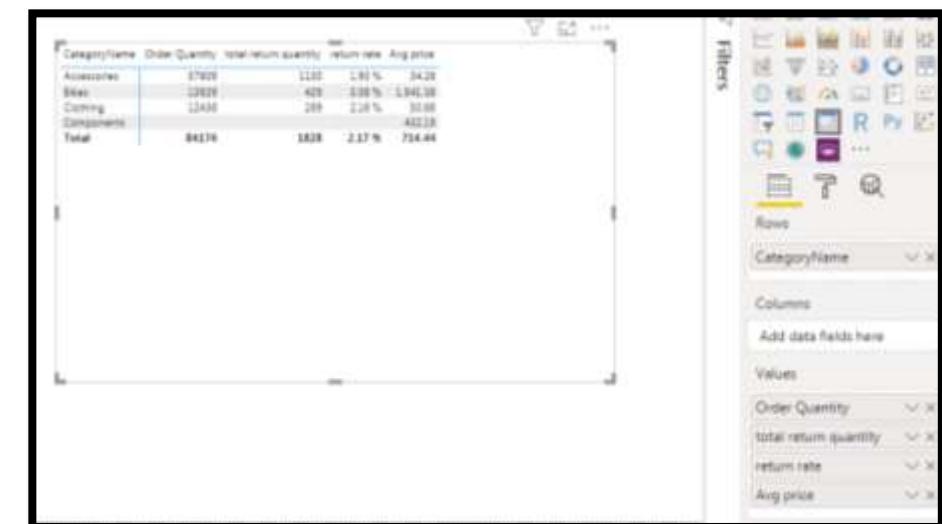
Demo : calculate the Avg of product price

- Go the report view
- Go to product table tab and create new measure.
- Write the followings DAX

```
Avg price = AVERAGE(AW_Product_Lookup[ProductPrice])
```

- Drag (Avg price) to the value field from the matrix
- Format the (price) of (Avg price) to currency
- Drag the (subcategory name)to the column

CategoryName	Order Quantity	Total return quantity	Return rate	Avg price
Accessories	87979	1130	1.25 %	34.28
Bikes	12929	429	3.38 %	354.58
Clothing	12438	289	2.31 %	30.00
Components	84628	4828	2.17 %	402.38
Total	384376	1828	2.17 %	314.44



Demo (Cont)

- Drag the product name to the (row) field of matrix

The screenshot shows a data visualization interface with a matrix view on the left and a filter sidebar on the right.

Matrix View:

CategoryName	Order Quantity	total quantity	return quantity	return rate	Avg price
Accessories	57809	1130	195	1.95 %	34.26
Bikes	13929	429	308	3.08 %	1,541.38
Clothing	12436	269	216	2.16 %	50.68
Components			432	19	432.19
Total	84174	1828	217	2.17 %	714.44

Filter Sidebar:

- Filters:** A grid of icons for different filter types.
- Rows:** A list of selected dimensions:
 - ProductName
 - SubcategoryName
 - CategoryName

Quiz

Q1.

- **Which two types of fields can DAX be used to create**

Calculated Columns & Calculated Fields

Calculated Columns & Measures

Calculated Fields & Measures

Measures & Calculated Items

- **What's the difference between calculated columns and measures**

Calculated columns understand row context

Measures understand filter context

Calculated column values are stored in tables

All of the above

-
- Which of the following formulas would make the *most* sense as a calculated column?

=SUM(Sales[quantity])

=AVERAGE(Products[RetailPrice])

=DISTINCTCOUNT(Customers[CustomerName])

=Products[RetailPrice] * 0.75

-
- TRUE or FALSE: The *Grand Total* cell in a Power BI visualization calculates by summing the measure values in the rows above it

True

False

- **When would you need to use the && operator?**

To concatenate two text strings

To create an AND condition between two expressions

To concatenate more than two text strings

To create an AND condition between more than two expressions

Exercise: Analyzing Data with DAX Calculations in Power BI

- Using the Adventure Works report, complete the following:
- **1)** In the **DATA** view, create the following **calculated columns**:
- In the **AW_Customer_Lookup** table, add a new column named "**Customer Priority**" that equals "*Priority*" for customers who are under 50 years old and have an annual income of greater than \$100,000, and "*Standard*" otherwise
- In the **AW_Product_Lookup** table, add a new column named "**Price Point**", based on the following criteria
 - *If the product price is greater than \$500, Price Point = "High"*
 - *If the product price is between \$100 and \$500, Price Point = "Mid-Range"*
 - *If the product price is less than or equal to \$100, Price Point = "Low"*
- In the **AW_Calendar_Lookup** table, add a new column named "**Short Day**" to extract and capitalize the first three letters from the **Day Name** column
- In the **AW_Product_Lookup** table, add a column named "**SKU Category**" to extract the first two characters from the **ProductSKU** field
 - **BONUS:** *Modify the SKU Category function to return any number of characters up to the first dash (Hint: You may need to "search" long and hard for that dash...)*

2) In the **REPORT** view, create the following **measures** (*Use a matrix visual to match the "spot check" values provided*)

- Create a measure named "**Product Models**" to calculate the number of unique product model names
- **Spot check:** You should see a total of **119** unique product models
- Create a measure named "**ALL Returns**" to calculate the grand total number of returns, regardless of the filter context
 - **Spot check:** You should see a total of **1,809** returns
- Create a measure to calculate "**% of All Returns**"
 - **Spot check:** You should see a value of **61.64%** for the **Accessories** product category
- Create a measure named "**Bike Returns**" to calculate total returns for bikes specifically
 - **Spot check:** You should see a total of **427** bike returns

-
- Create a measure named "**Total Cost**", by multiplying order quantities by product costs at the row-level
 - *Spot check: You should see a total cost of \$14,456,986.32*
 - Once you've calculated **Total Cost**, create a new measure for "**Total Profit**", defined as the total revenue minus the total cost
 - *Spot check: You should see a total profit of \$10,457,580.86*

-
- Total Cost = sumx('AW-Sales', 'AW-Sales'[OrderQuantity] * related(AW_Product_Lookup[ProductCost]))
 - total profit = [total revenue] - [Total Cost]